




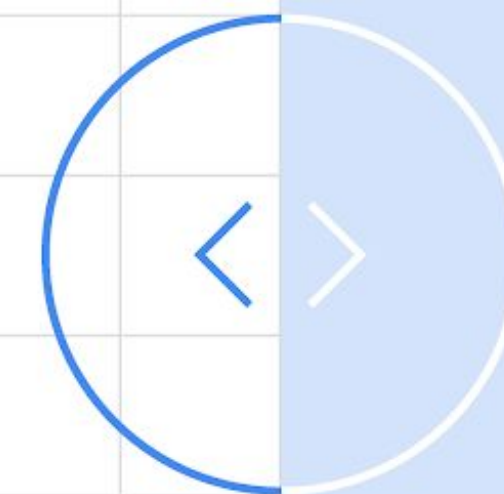
Hello, TensorFlow :)

Basic Deep Learning recipes with TensorFlow




Sayak Paul
PyImageSearch
[@RisingSayak](#)

Google Developers



Acknowledgement


- The entire PyImageSearch team
- Laurence Moroney (Google)
- François Chollet (Google)
- Soonson Kwon (Google)



If terms like TensorFlow and Deep Learning scare you,
don't worry you're not alone!

Prerequisites:

- Python (at least 6 months)
- Basic high school math



Agenda

- What is TensorFlow?
 - Tensor
 - Flow
- An introduction to Deep Learning
 - What is *learning* here?
 - Why the word *deep*?
- Let's get coding!
- QA (if time permits)

What is TensorFlow?

As per [tensorflow.org](https://www.tensorflow.org) -

“The core open source library to help you develop and train ML models.”

What is TensorFlow?

But at its core, *TensorFlow is a library for doing numerical computation.*

Why is it called TensorFlow?

Two components in there -

- Tensor
- Flow

Tensors

Multi-dimensional arrays!

		<i>Columns</i> →				
		0	1	2	3	4
<i>Rows</i> ↓	0	5	12	17	9	3
	1	13	4	8	14	1
	2	9	6	3	7	21

2D Array of size 3 x 5

Tensors

You can call it a container of multi-dimensional arrays as well.

```
>>> X_train.shape
(60000, 28, 28)

>>> X_train[0].shape
(28, 28)
```

Flow

Flow represents a *directed computational graph*!

Flow

Flow represents a *directed computational graph*!



Can you represent the following equation in terms of a graph?

$$y = wX + b$$

Flow

Flow represents a *directed computational graph*!

Yes, you *can*!



Flow

Flow represents a *computational graph*!

Long story cut short, TensorFlow represents all the variables and the operations in terms of a computation graph.*

*In *eager mode*, it *does not* construct a computation graph beforehand.

The promise of Machine Learning




Source: TensorFlow team

The promise of Machine Learning

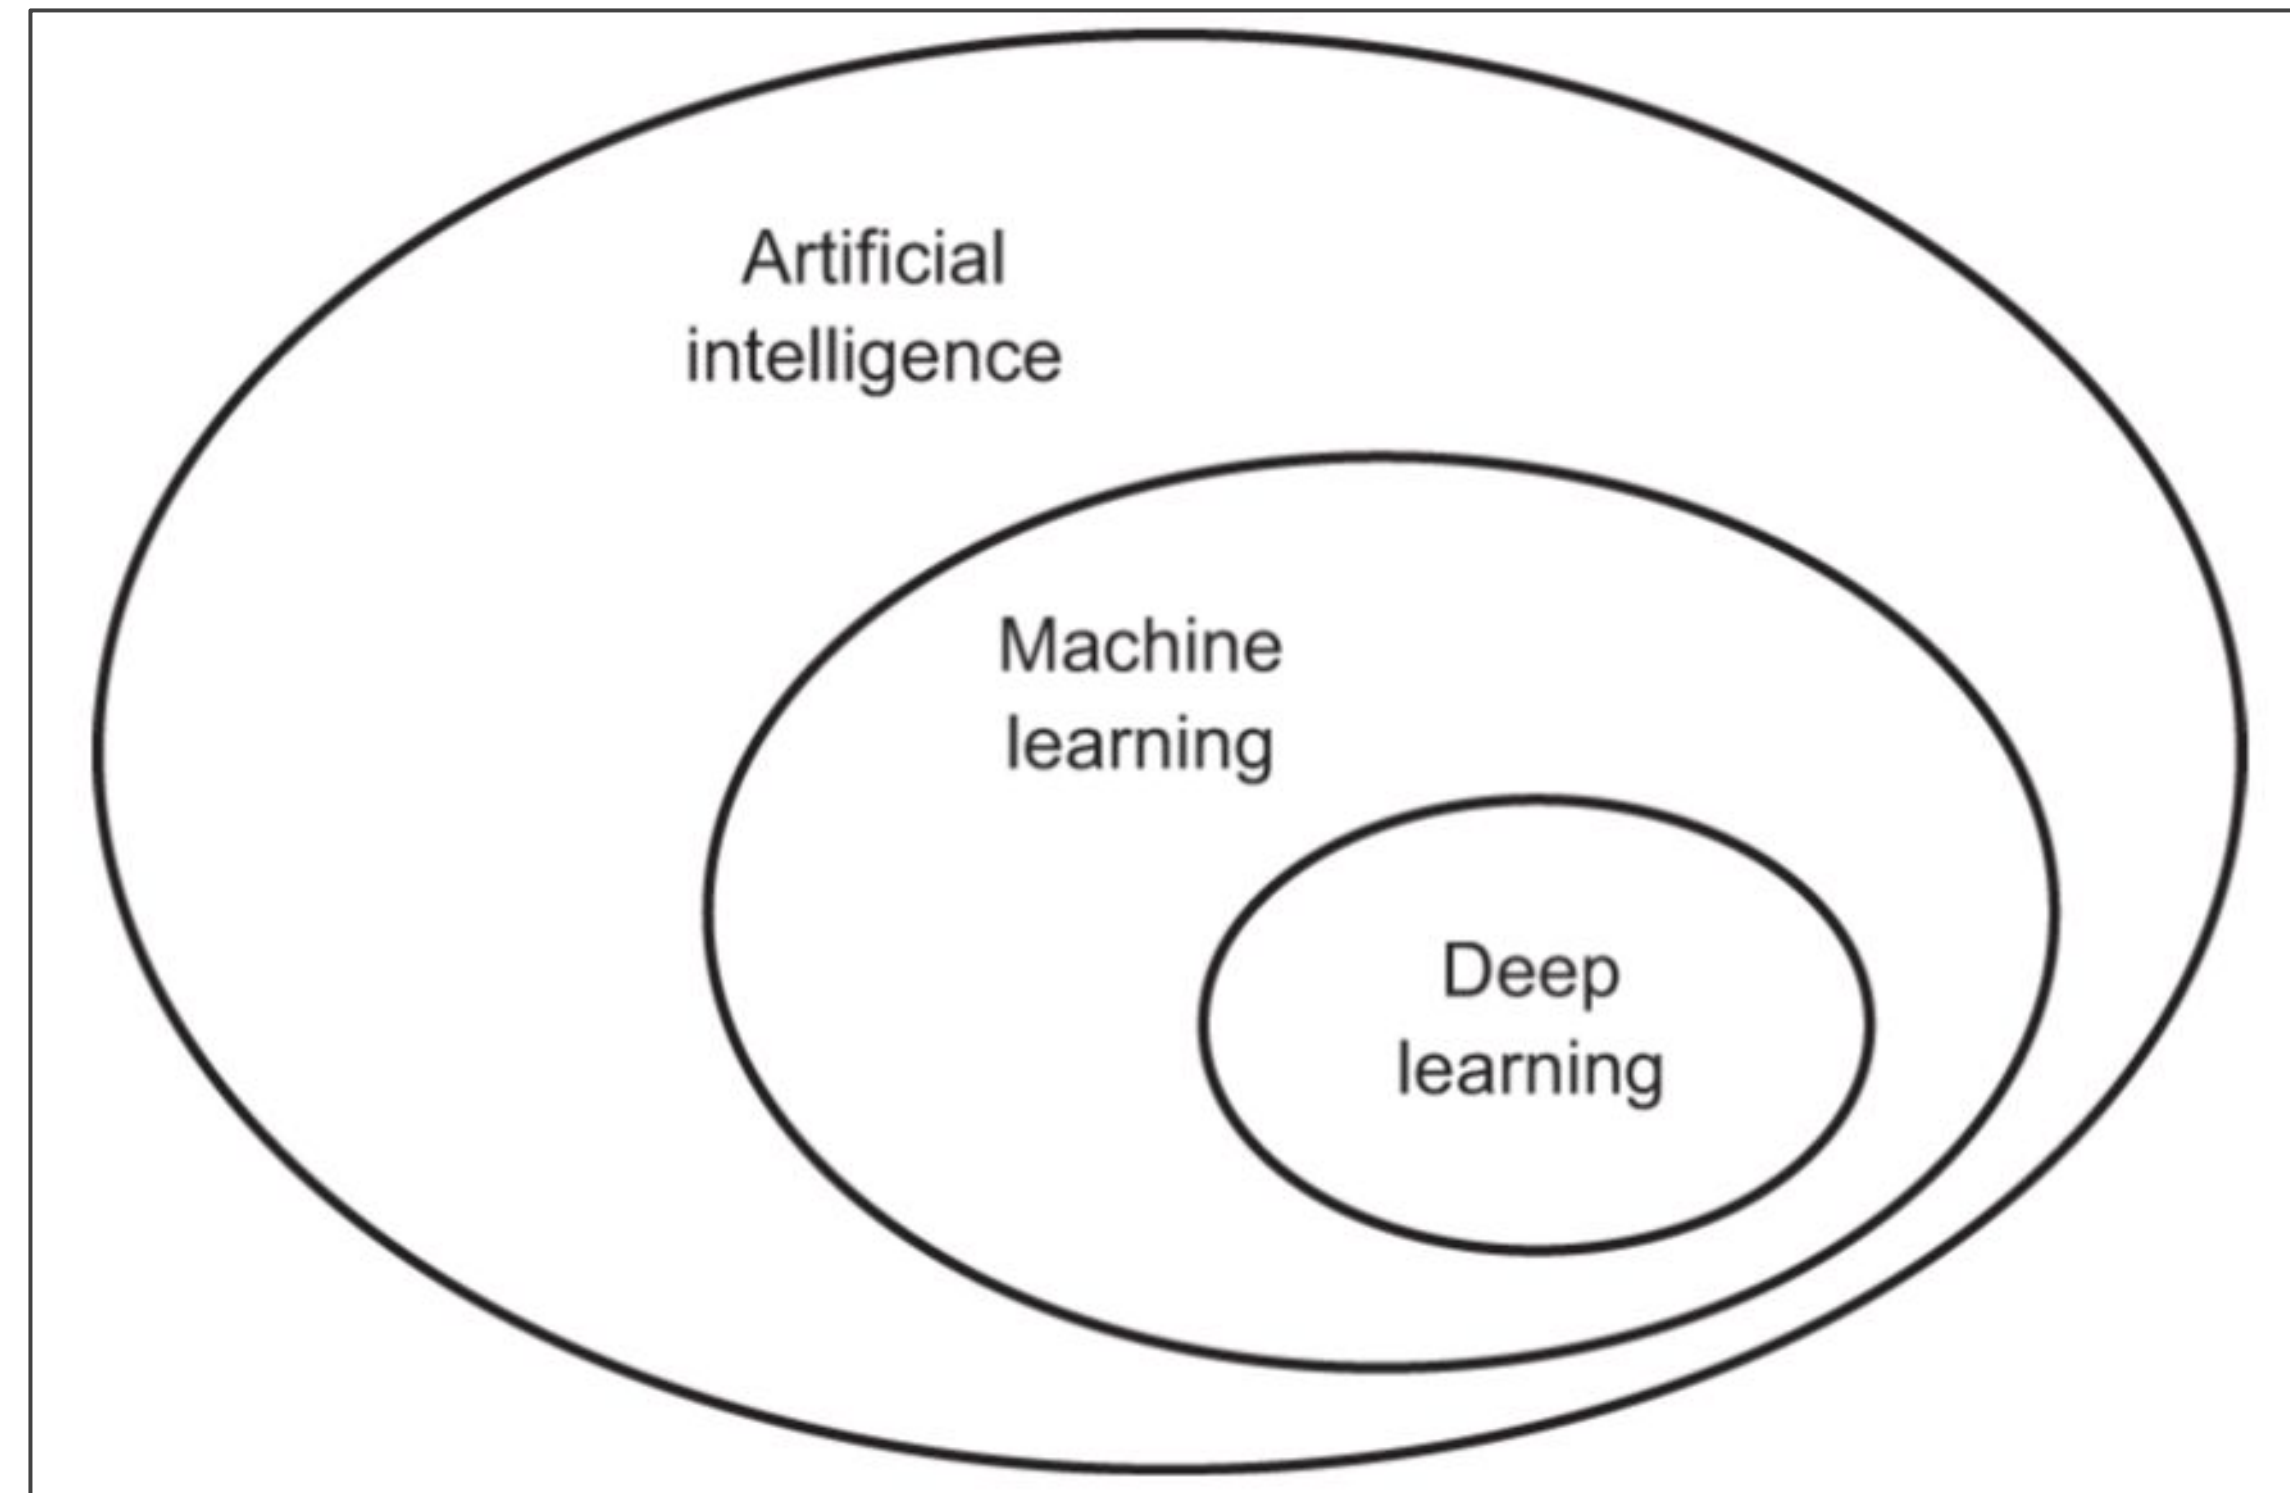


Source: TensorFlow team


Aren't we doing Deep Learning today?



Yes, of course!



Source: [Deep Learning with Python](#)




So, what conclusion can we draw from that figure?

Going back to this figure ...



Source: TensorFlow team

Why this might be **useful**?



Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```


Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```

```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```

What now?

Our daily (digital) life

- Google Photos
- Next word suggestion while you're typing
- Tag suggestions in Facebook photos
- ...

Two fundamental elements

- Data
- Labels (Answers)

An example

age	secondary makrs	higher secondary marks	admission
20	356	471	0
18	421	405	1
19	321	300	0
18	423	400	1
...

An example

age	secondary makrs	higher secondary marks	admission
20	356	471	0
18	421	405	1
19	321	300	0
18	423	400	1
...

Multiple rows and multiple columns!

An example

age	secondary makrs	higher secondary marks	admission
20	356	471	0
18	421	405	1
19	321	300	0
18	423	400	1
...

Rows: **Data points/Instances**
Columns: **Features**


An example

age	secondary makrs	higher secondary marks	admission
20	356	471	0
18	421	405	1
19	321	300	0
18	423	400	1
...
Data			Labels/Answers

Rows: **Data points/Instances**
Columns: **Features (the last one is for *Labels*)**

Task at hand

How can we separate the following data points?



Task at hand

To separate the data points we need to know the ***underlying pattern of the data points*** that separates them.


Task at hand

To separate the data points we need to know the ***underlying pattern of the data points*** that separates them.

You cannot distinguish between a cat and a dog if you don't know about them.


Task at hand

Quick thought: **The data points are linearly separable**



Task at hand

But how do we find that line?



Meet your new fellas – parameters

- Weights
- Biases

Meet your new fellas – parameters

- We apply a series of **transformations** on our input data (the figure) with weights and biases.

Meet your new fellas – parameters

- We apply a series of **transformations** on our input data (the figure) with weights and biases.
- These transformations are needed to uncrumple our input data.

Meet your new fellas - parameters

You have already seen how these transformations look like -

$$y = wX + b$$

Meet your new fellas - parameters

You have already seen how these transformations look like -

$$y = wX + b$$

X = Data points

W = Weights

B = Bias

“Talk is cheap. Show me the code!”

Head over to this link 🖱️ bit.ly/ht-colab

- Transform our input data
- Calculate error
- Derivative 😓
- And more ...

From the code


The way we minimized the loss function (Mean Squared Error) comprises of the following -

- Linear regression
- Gradient-based learning
- Backpropagation

Linear Regression

The entire system that we built to minimize our loss function -

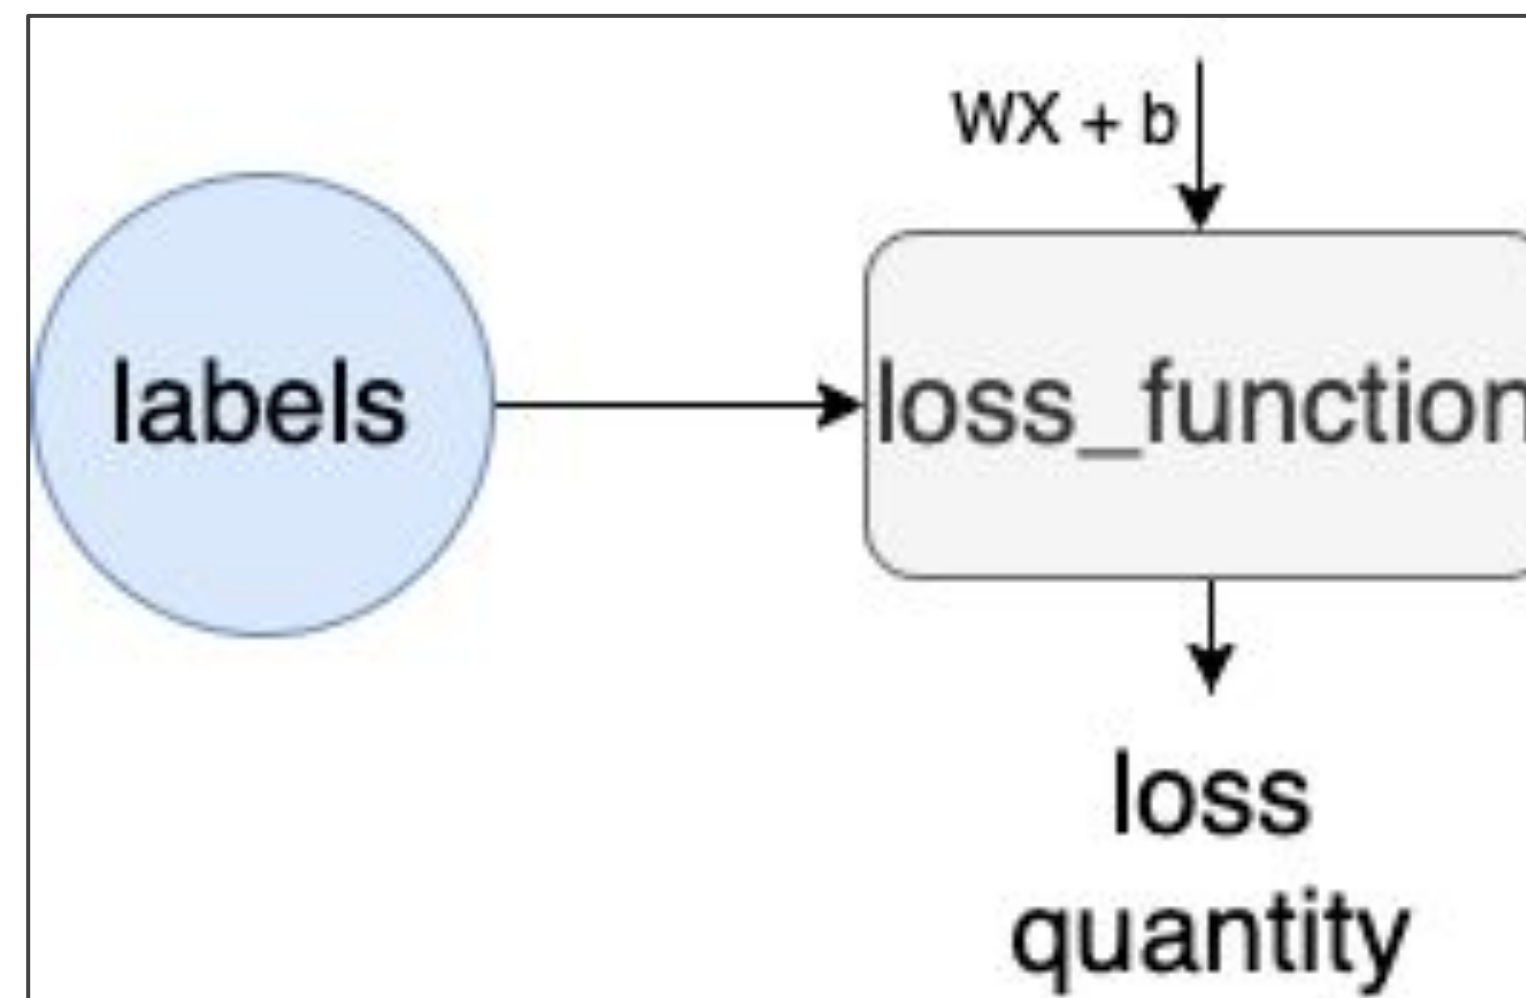
a. We applied transformations to our input data.



Linear Regression

The entire system that we built to minimize our loss function -

- b. We calculated the loss between the results of those transformations and our labels.

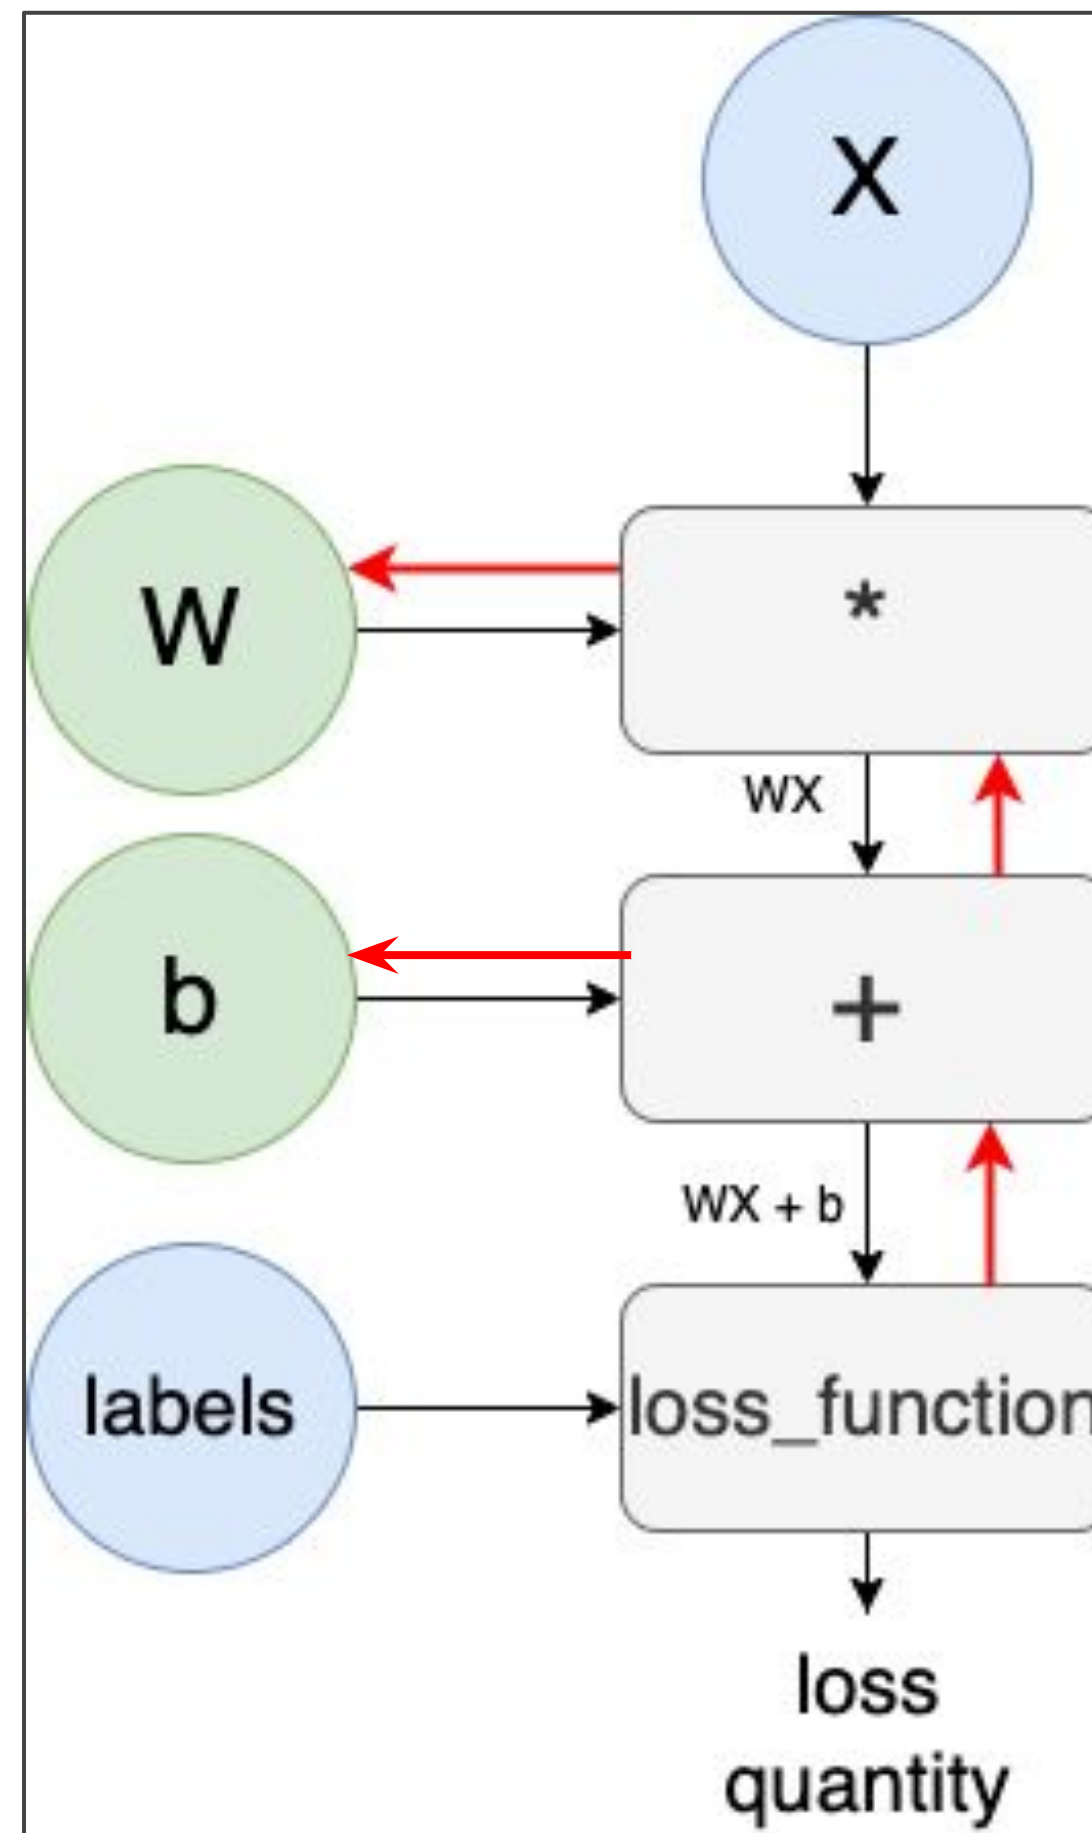


Linear Regression

The entire system that we built to minimize our loss function -

c. We then *calculated derivatives of the loss function* w.r.t the parameters
and *updated the parameters*.

Linear Regression

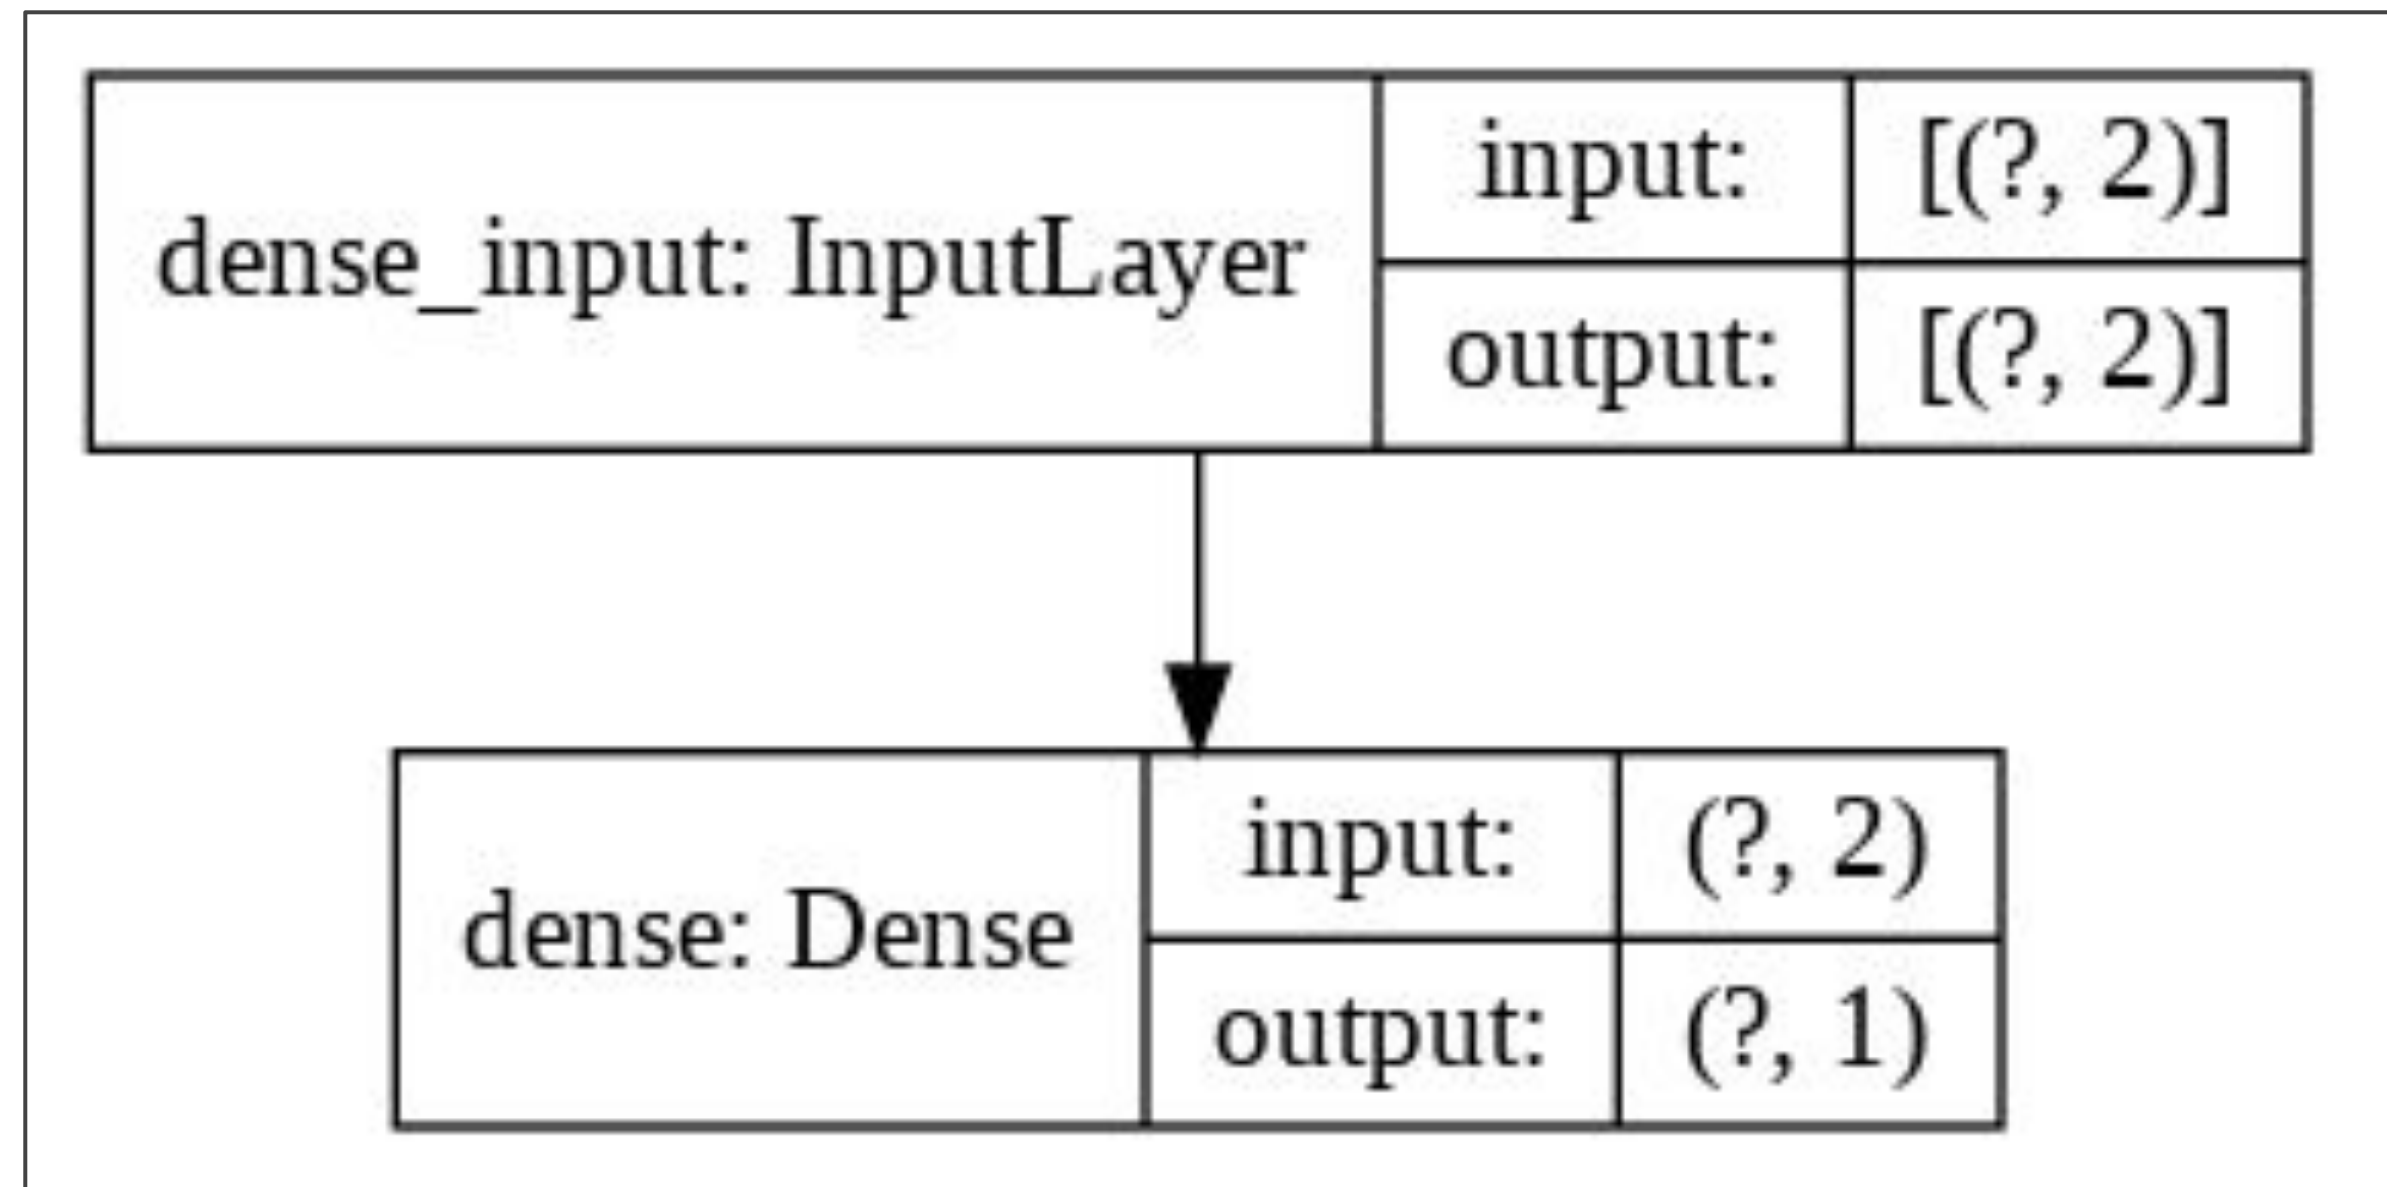


Linear Regression

The entire system that we built to minimize our loss function -

- a. We applied transformations to our input data.
- b. We calculated the loss between the results of those transformations and our labels.
- c. We then *calculated derivatives of the loss function w.r.t the parameters and updated the parameters.*
- d. We repeated **a**, **b**, and **c**.

Schematically our system is



Gradient-based *learning*

```
def minimize_loss(x, y):  
    with tf.GradientTape() as tape:  
        # Transform data and compute the loss  
        predictions = compute_predictions(x)  
        loss = compute_loss(y, predictions)  
  
        # Compute the derivative/gradients  
        dloss_dw, dloss_db = tape.gradient(loss, [w, b])  
  
    # Update the parameters and return loss  
    w.assign_sub(LEARNING_RATE * dloss_dw)  
    b.assign_sub(LEARNING_RATE * dloss_db)  
    return loss
```

Gradient-based *learning*

- Learning, wait what?
- *Learning* is a fancy term for **function approximation**.
- In our case, we tried to *find the underlying pattern* of the input data points *that separates them one another*.

Gradient-based *learning*

- Learning, wait what?
- *Learning* is a fancy term for **function approximation**.
- In our case, we tried to *find the underlying pattern* of the input data points *that separates them one another*.

These underlying patterns are nothing but functions and **finding these patterns = function approximation**

Gradient-based *learning*

- Remember we minimized the loss function?
- In other words, we ***trained*** a system to minimize the loss function.
- We did so with the help of ***gradient-based learning***.
- As we *updated the parameters* in this process, they are also called ***trainable variables/learnable parameters***.

Gradient-based *learning*

- Remember we minimized the loss function?
- In other words, we ***trained*** a system to minimize the loss function.
- We did so with the help of ***gradient-based learning***.
- As we *updated the parameters* in this process, they are also called ***trainable variables/learnable parameters***.

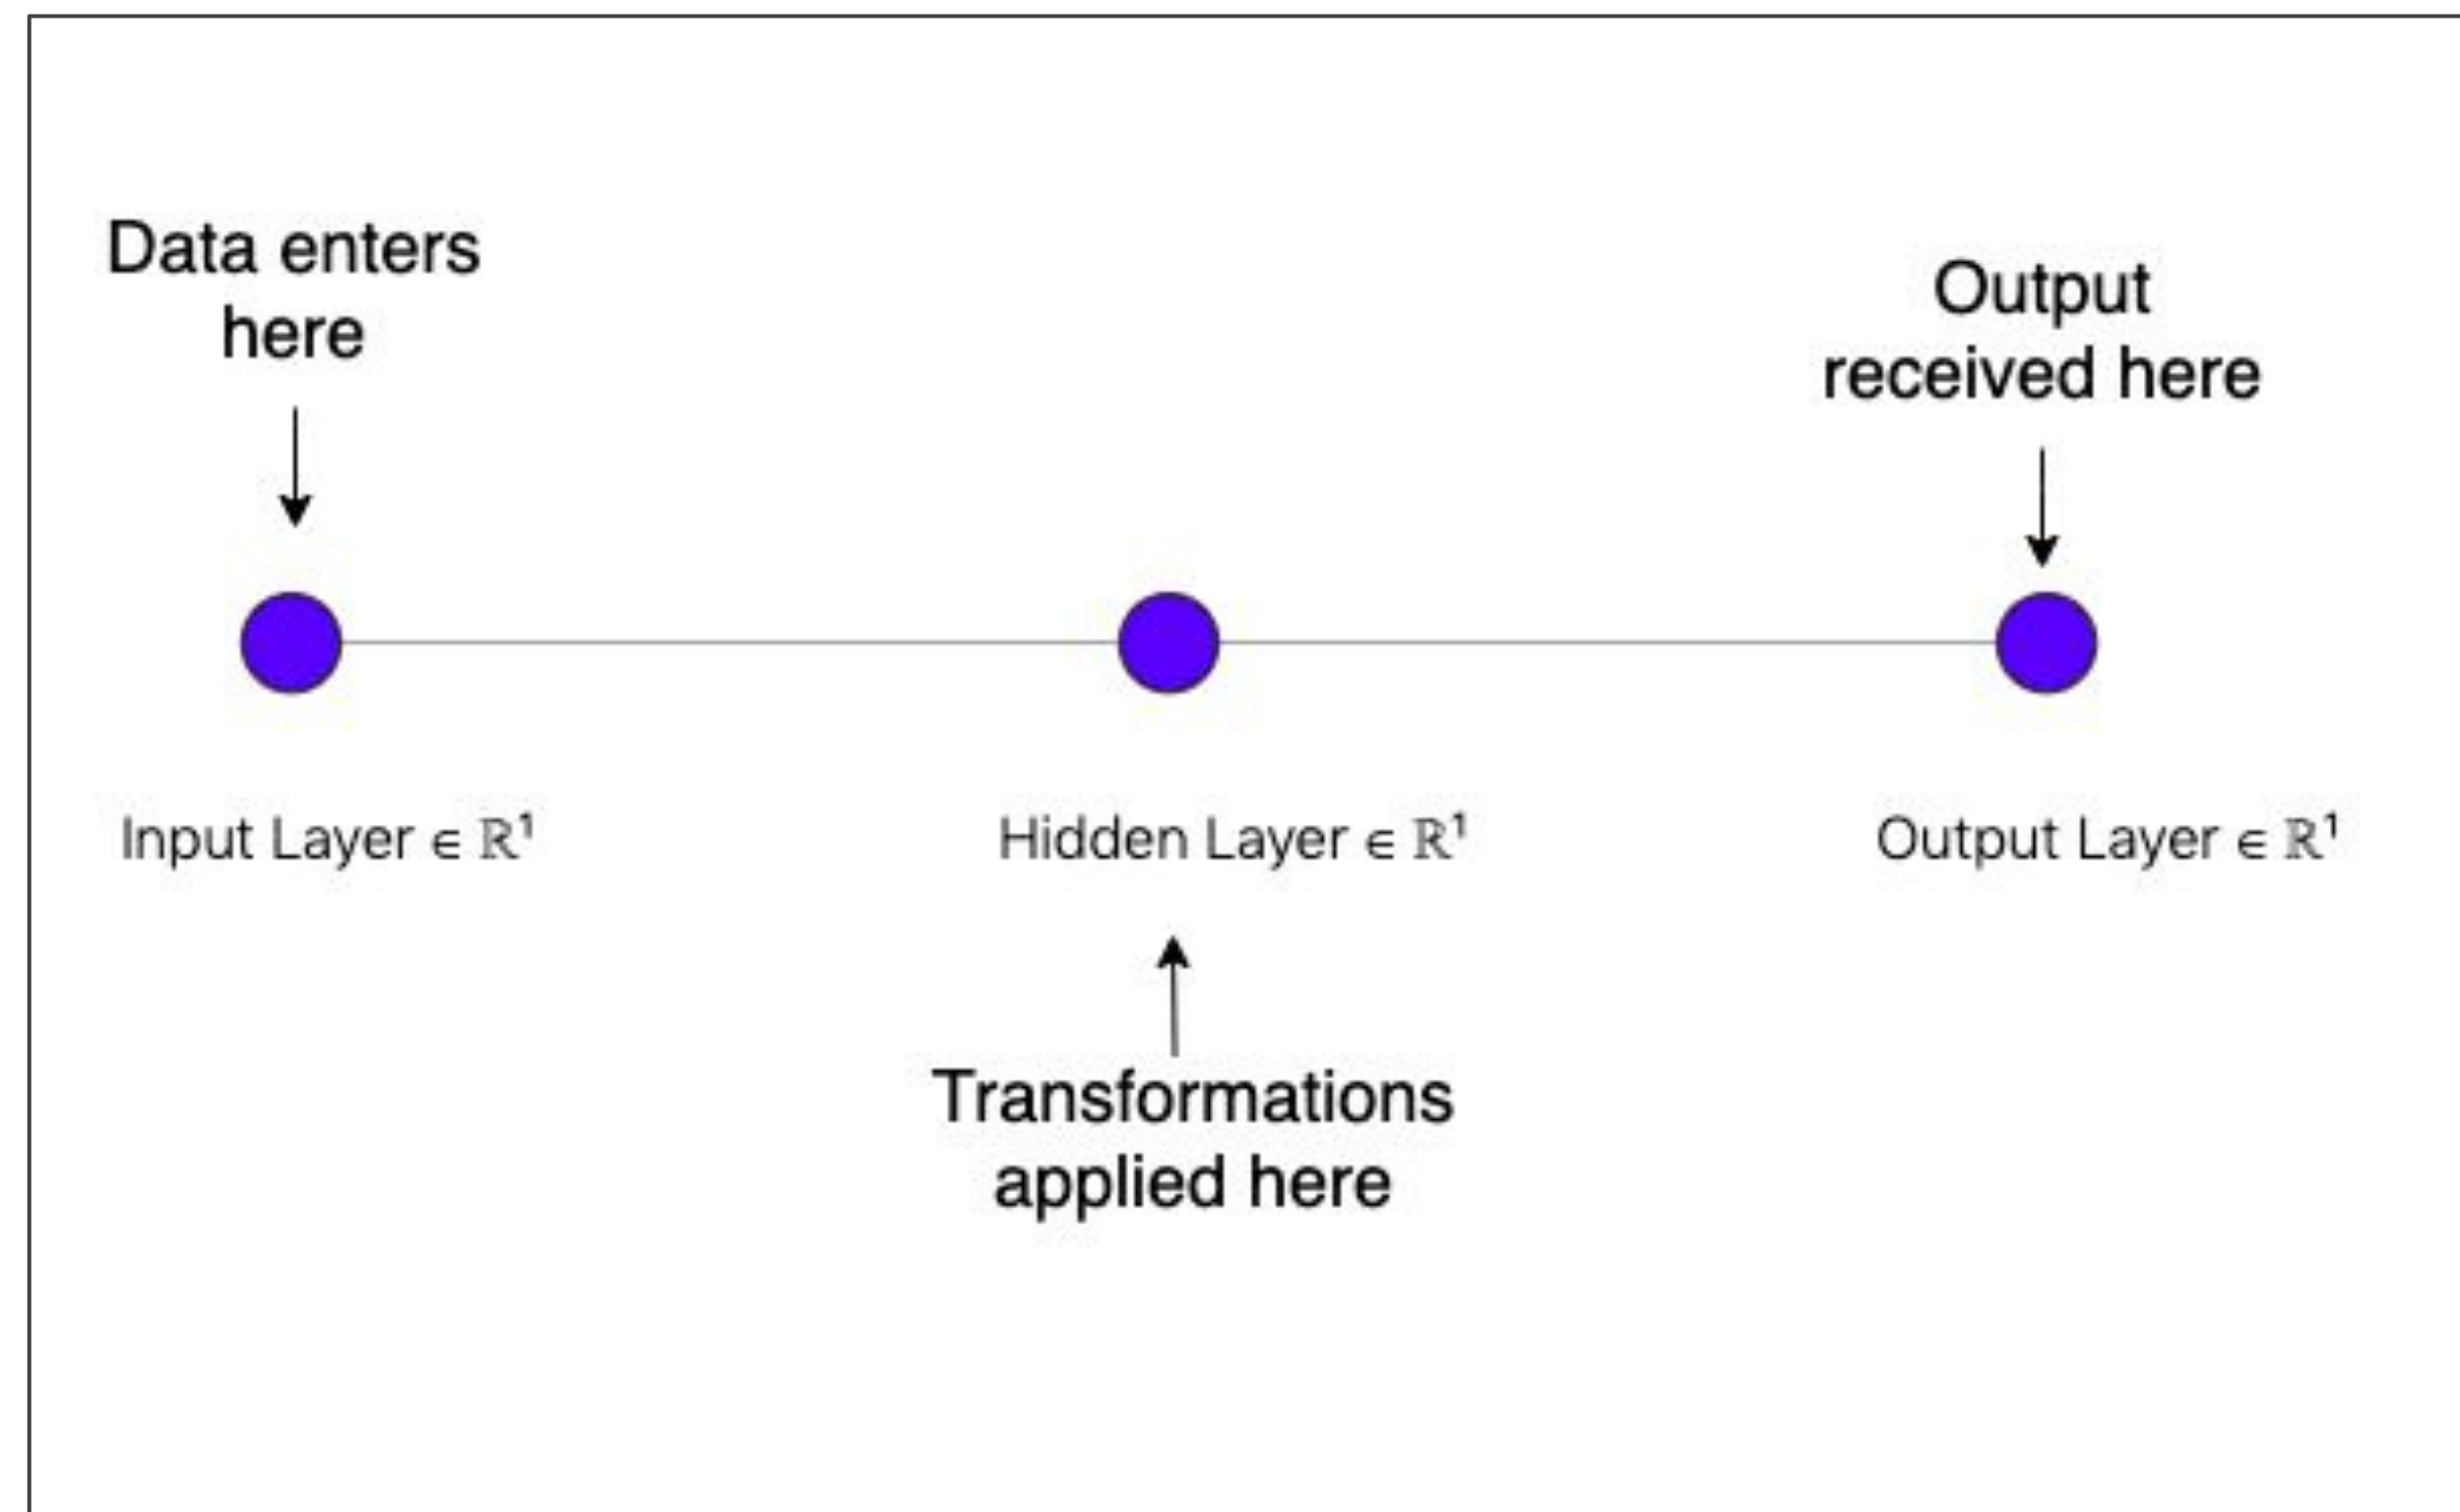
This form of learning/training is also known as ***Supervised Learning***.

Backpropagation

```
def minimize_loss(x, y):  
    with tf.GradientTape() as tape:  
        # Transform data and compute the loss  
        predictions = compute_predictions(x)  
        loss = compute_loss(y, predictions)  
  
        # Compute the derivative/gradients  
        dloss_dw, dloss_db = tape.gradient(loss, [w, b])  
  
        # Update the parameters and return loss  
        w.assign_sub(LEARNING_RATE * dloss_dw)  
        b.assign_sub(LEARNING_RATE * dloss_db)  
    return loss
```



Meaning of *deep*

Another diagram of our system could be -



Meaning of *deep*

What if a diagram looked like so?



Meaning of *deep*

- We can add *depth* to our system in a similar manner.

Meaning of *deep*

- We can add *depth* to our system in a similar manner.
- A quick question –

Did you find any similarity in between the previous figure and the ***neurons*** inside our brain?


Meaning of *deep*

- We can add *depth* to our system in a similar manner.
- A quick question –

Did you find any similarity in between the previous figure and the *neurons* inside our brain?

That is how the powerful **neural networks** look like

This is no way even the
scratch of the surface!



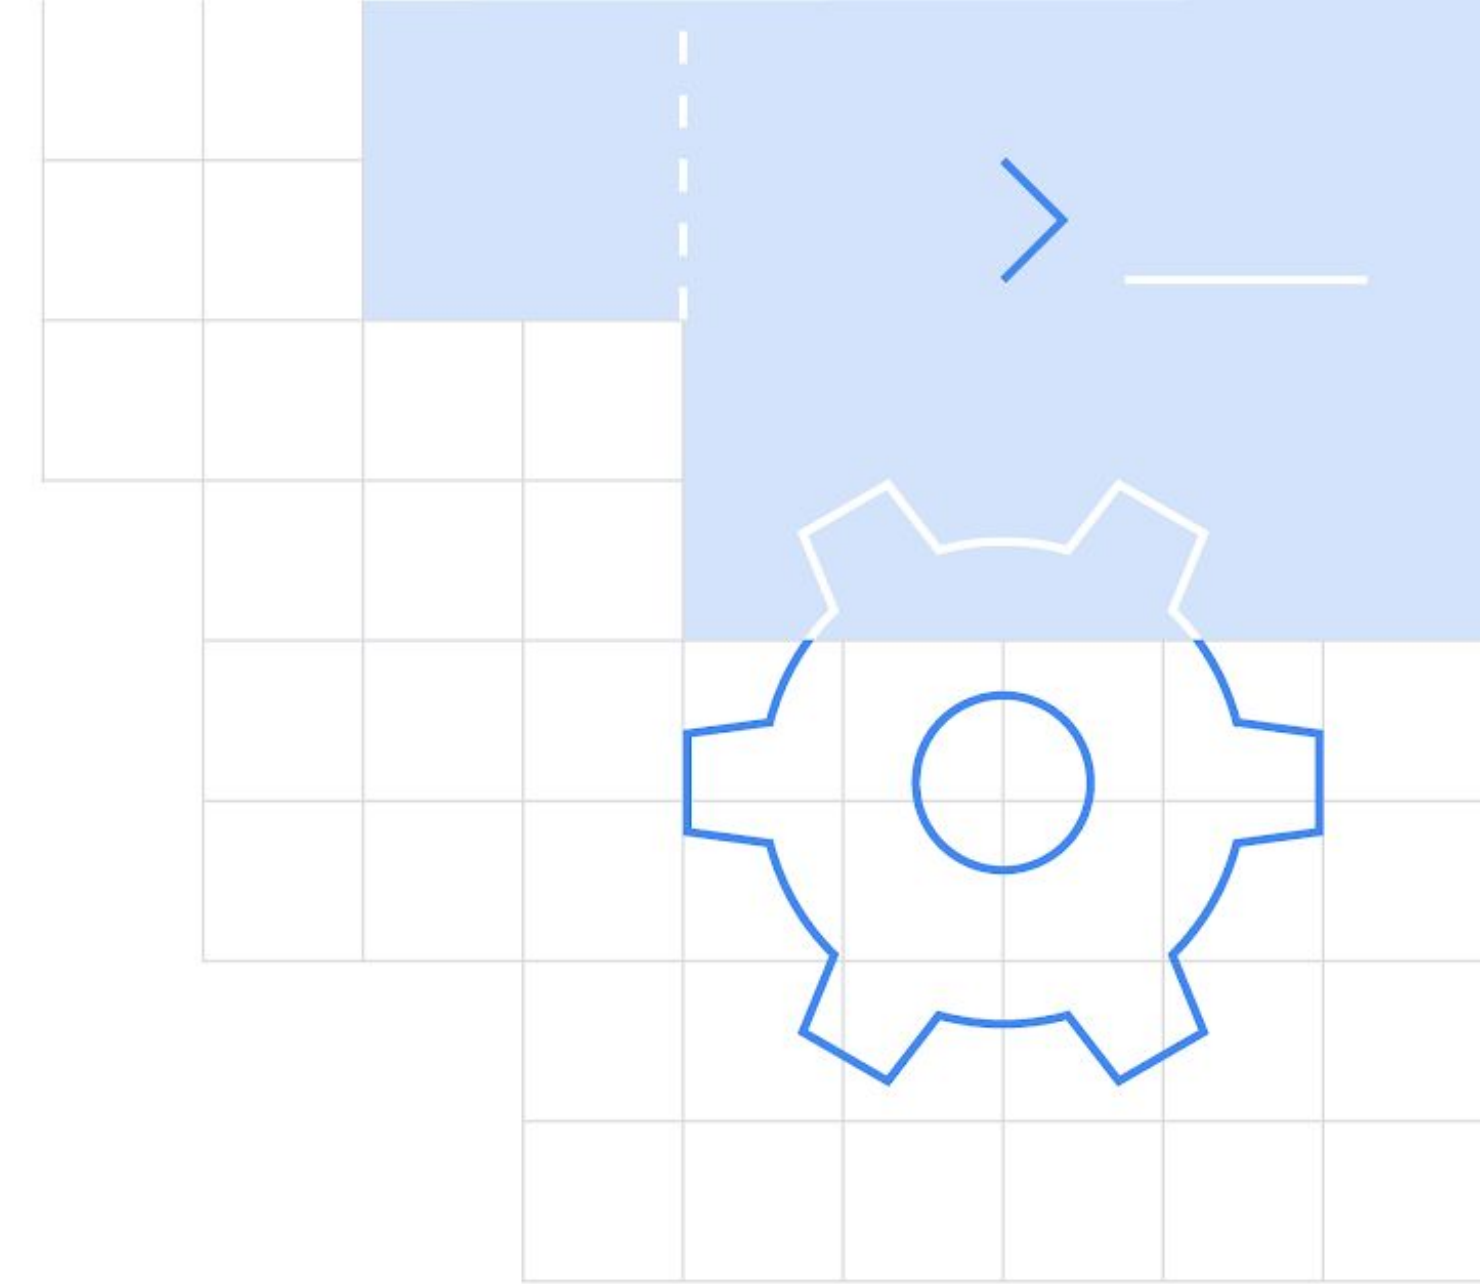
If this seemed interesting

Consider the following resource to jump in this ocean -

- [Deep Learning With Python](#) by **François Chollet**
- [TensorFlow in Practice Specialization](#) by **deeplearning.ai & Laurence Moroney**
- [Neural Networks and Deep Learning](#) by **Michael Nielsen**
- [TensorFlow, Keras and deep learning, without a PhD](#) by **Martin Gorner**
- Cannot fit more in one single slide, sorry :(

Slides available here -

bit.ly/gdg-goat-20



Thank You!



Sayak Paul
PyImageSearch
[@RisingSayak](#)

