

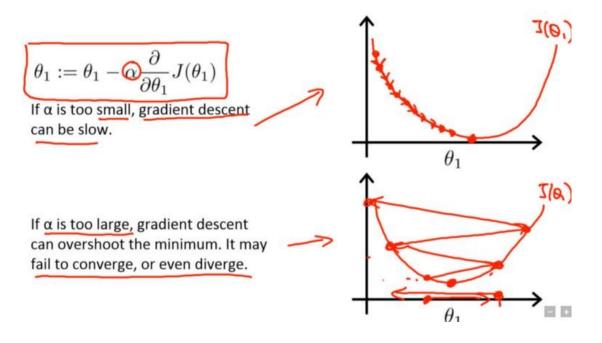
Introduction to Cyclical Learning Rates for training Neural Nets

Sayak Paul Project Instructor (Google DevFest, Kolkata, India)



- Why are *learning rates* used?
- Some existing approaches for choosing the right learning rate
- What are the shortcomings of these approaches?
- Need of a systematic approach for setting the learning rate –
 Cyclical Learning Rates (CLR)
- What is CLR?
- Some amazing results shown by CLR
- Conclusion

Learning is an important *hyperparameter* for adjusting the weights of a network with respect to the loss gradient.

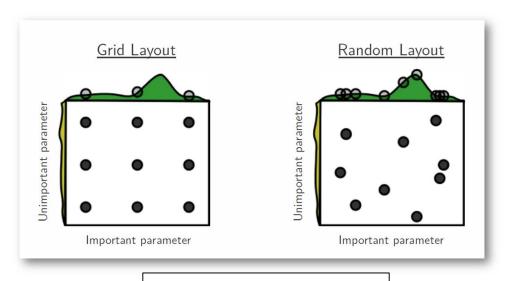


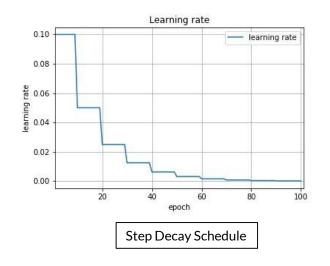
Source: Andrew Ng's lecture notes from Coursera

Some of the existing approaches for choosing the right



- Trying out different learning rates for a problem.
- Grid-searching/Random-searching.
- Adaptive Learning Rates / Learning Rate Schedules.





Grid and Random layout of parameters

Problems with these approaches



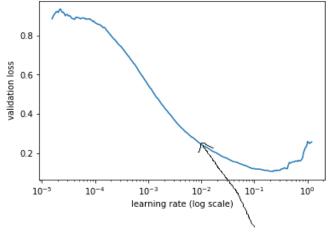
- Computationally costly.
- Gives no early clue if at all the result would get better.

Cyclical Learning Rates*



Proposed by Leslie N. Smith in his paper entitled "Cyclical Learning Rates for Training Neural Networks" in 2015.

• The idea is to simply keep increasing the learning rate from a very small value, until the loss stops decreasing.



The sweet spot!

<u>Source</u>

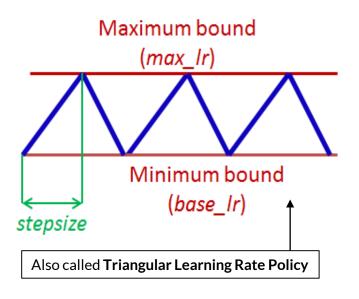
^{*} Cyclical Learning Rates for Training Neural Networks - Leslie N. Smith

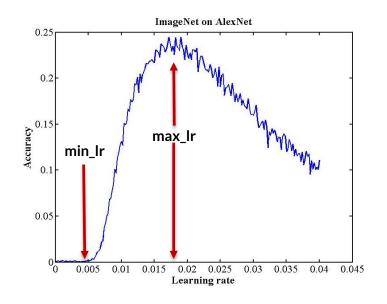
How are Cyclical Learning Rates (CLR) systematic?



- The main idea behind CLR is varying learning rates between min and max values.
- LR_Range_Test() is conducted for fixing the min and max values of learning rate.

• One step of increasing learning rate.

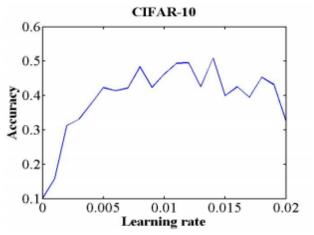




Choosing max_Ir and min_Ir

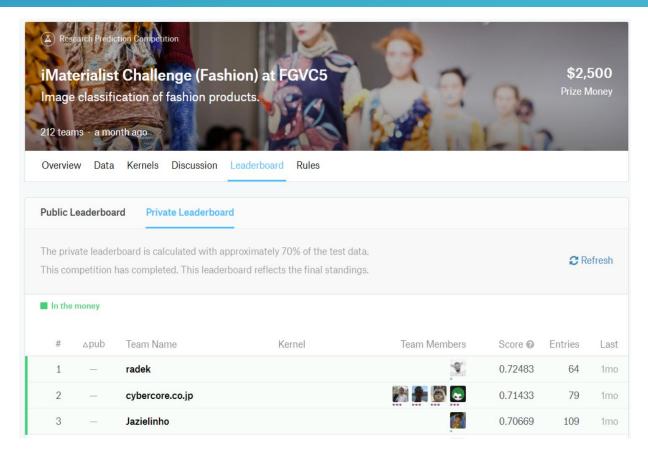


- Run the model for several epochs while letting the learning rate increase linearly (use triangular learning rate policy) between low and high learning rate values.
- Next, plot the accuracy versus learning rate curve.
- Note the learning rate value when the accuracy starts to increase and when the accuracy slows, becomes ragged, or starts to fall. These two learning rates are good choices for defining the range of the learning rates.



Some amazing results shown by CLR



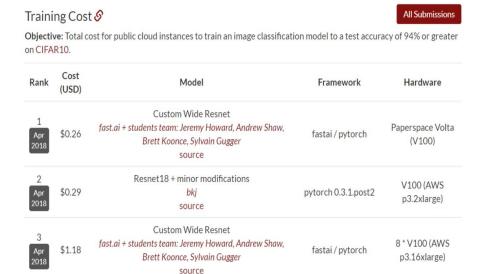


Some amazing results shown by CLR (contd.)



Image Classification on CIFAR10

Training Time § All Submissions Objective: Time taken to train an image classification model to a test accuracy of 94% or greater on CIFAR10. Time to 94% Rank Model Framework Hardware Accuracy Custom Wide Resnet fast.ai + students team: Jeremy Howard, Andrew 8 * V100 (AWS 0:02:54 fastai / pytorch Shaw, Brett Koonce, Sylvain Gugger p3.16xlarge) source Resnet18 + minor modifications pytorch V100 (AWS 0:05:41 bkj p3.2xlarge) 0.3.1.post2 source Custom Wide Resnet fast.ai + students team: Jeremy Howard, Andrew Paperspace Volta 0:06:45 fastai/pytorch (V100) Shaw, Brett Koonce, Sylvain Gugger source



DAWNBench Challenge Leaderboard and Leader's specs

Limitations of CLR



- Limited applicability.
- Seems to work only for Cifar-10 and resnets.
- But definitely provides a more systematic way for choosing learning rate than the earlier approaches.

Notable byproducts of CLR



- Learning rate annealing (SDGR).
- Differential Learning Rates.



Cyclical Learning Rates for Training Neural Networks

Leslie N. Smith U.S. Naval Research Laboratory, Code 5514 4555 Overlook Ave., SW., Washington, D.C. 20375

leslie.smith@nrl.navy.mil

Original CLR Paper



Introduction to Cyclical Learning Rates

Learn what cyclical learning rate policy is and how it can improve the training of a neural network.

DataCamp tutorial covering CLR

sayakpaul Merge pull request #3 from sayakpaul/sayakpaul-patch-3		Latest commit 61fc822 9 hours ago
☐ Introduction to Cyclical Learning Rates for training Neural Nets.pdf	Minor fixes	9 hours ago
LICENSE	Initial commit	5 days ago
README.md	Update README.md	5 days ago

Slides available on my Github (Username: sayakpaul)



Thank you!

Sayak Paul Project Instructor spsayakpaul@gmail.com