





Gotchas of transfer learning for image classification



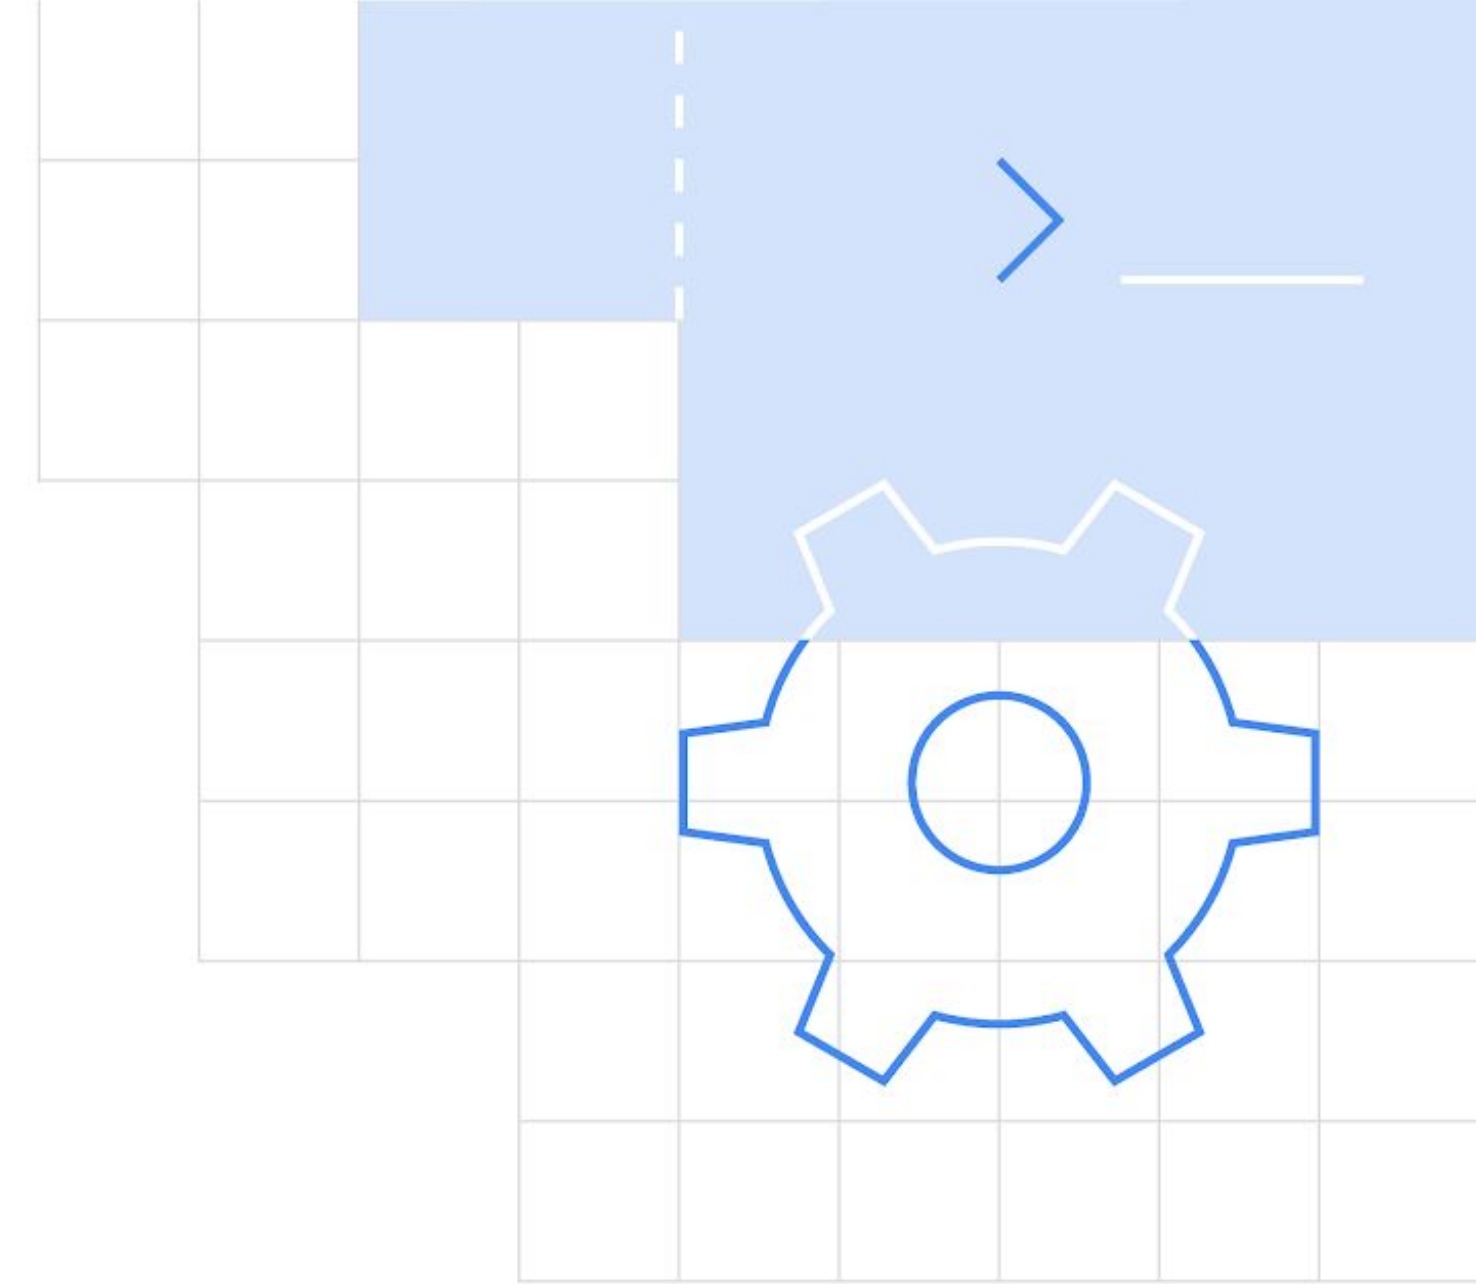
Sayak Paul
PyImageSearch
[@RisingSayak](#)

Google Developers



Ideal audience

- ML Developers that have trained an image classifier



Agenda

- Brief introduction to transfer learning
- The ImageNet moment
- Transfer learning for image classification
- Things to keep in mind while doing transfer learning
- QA

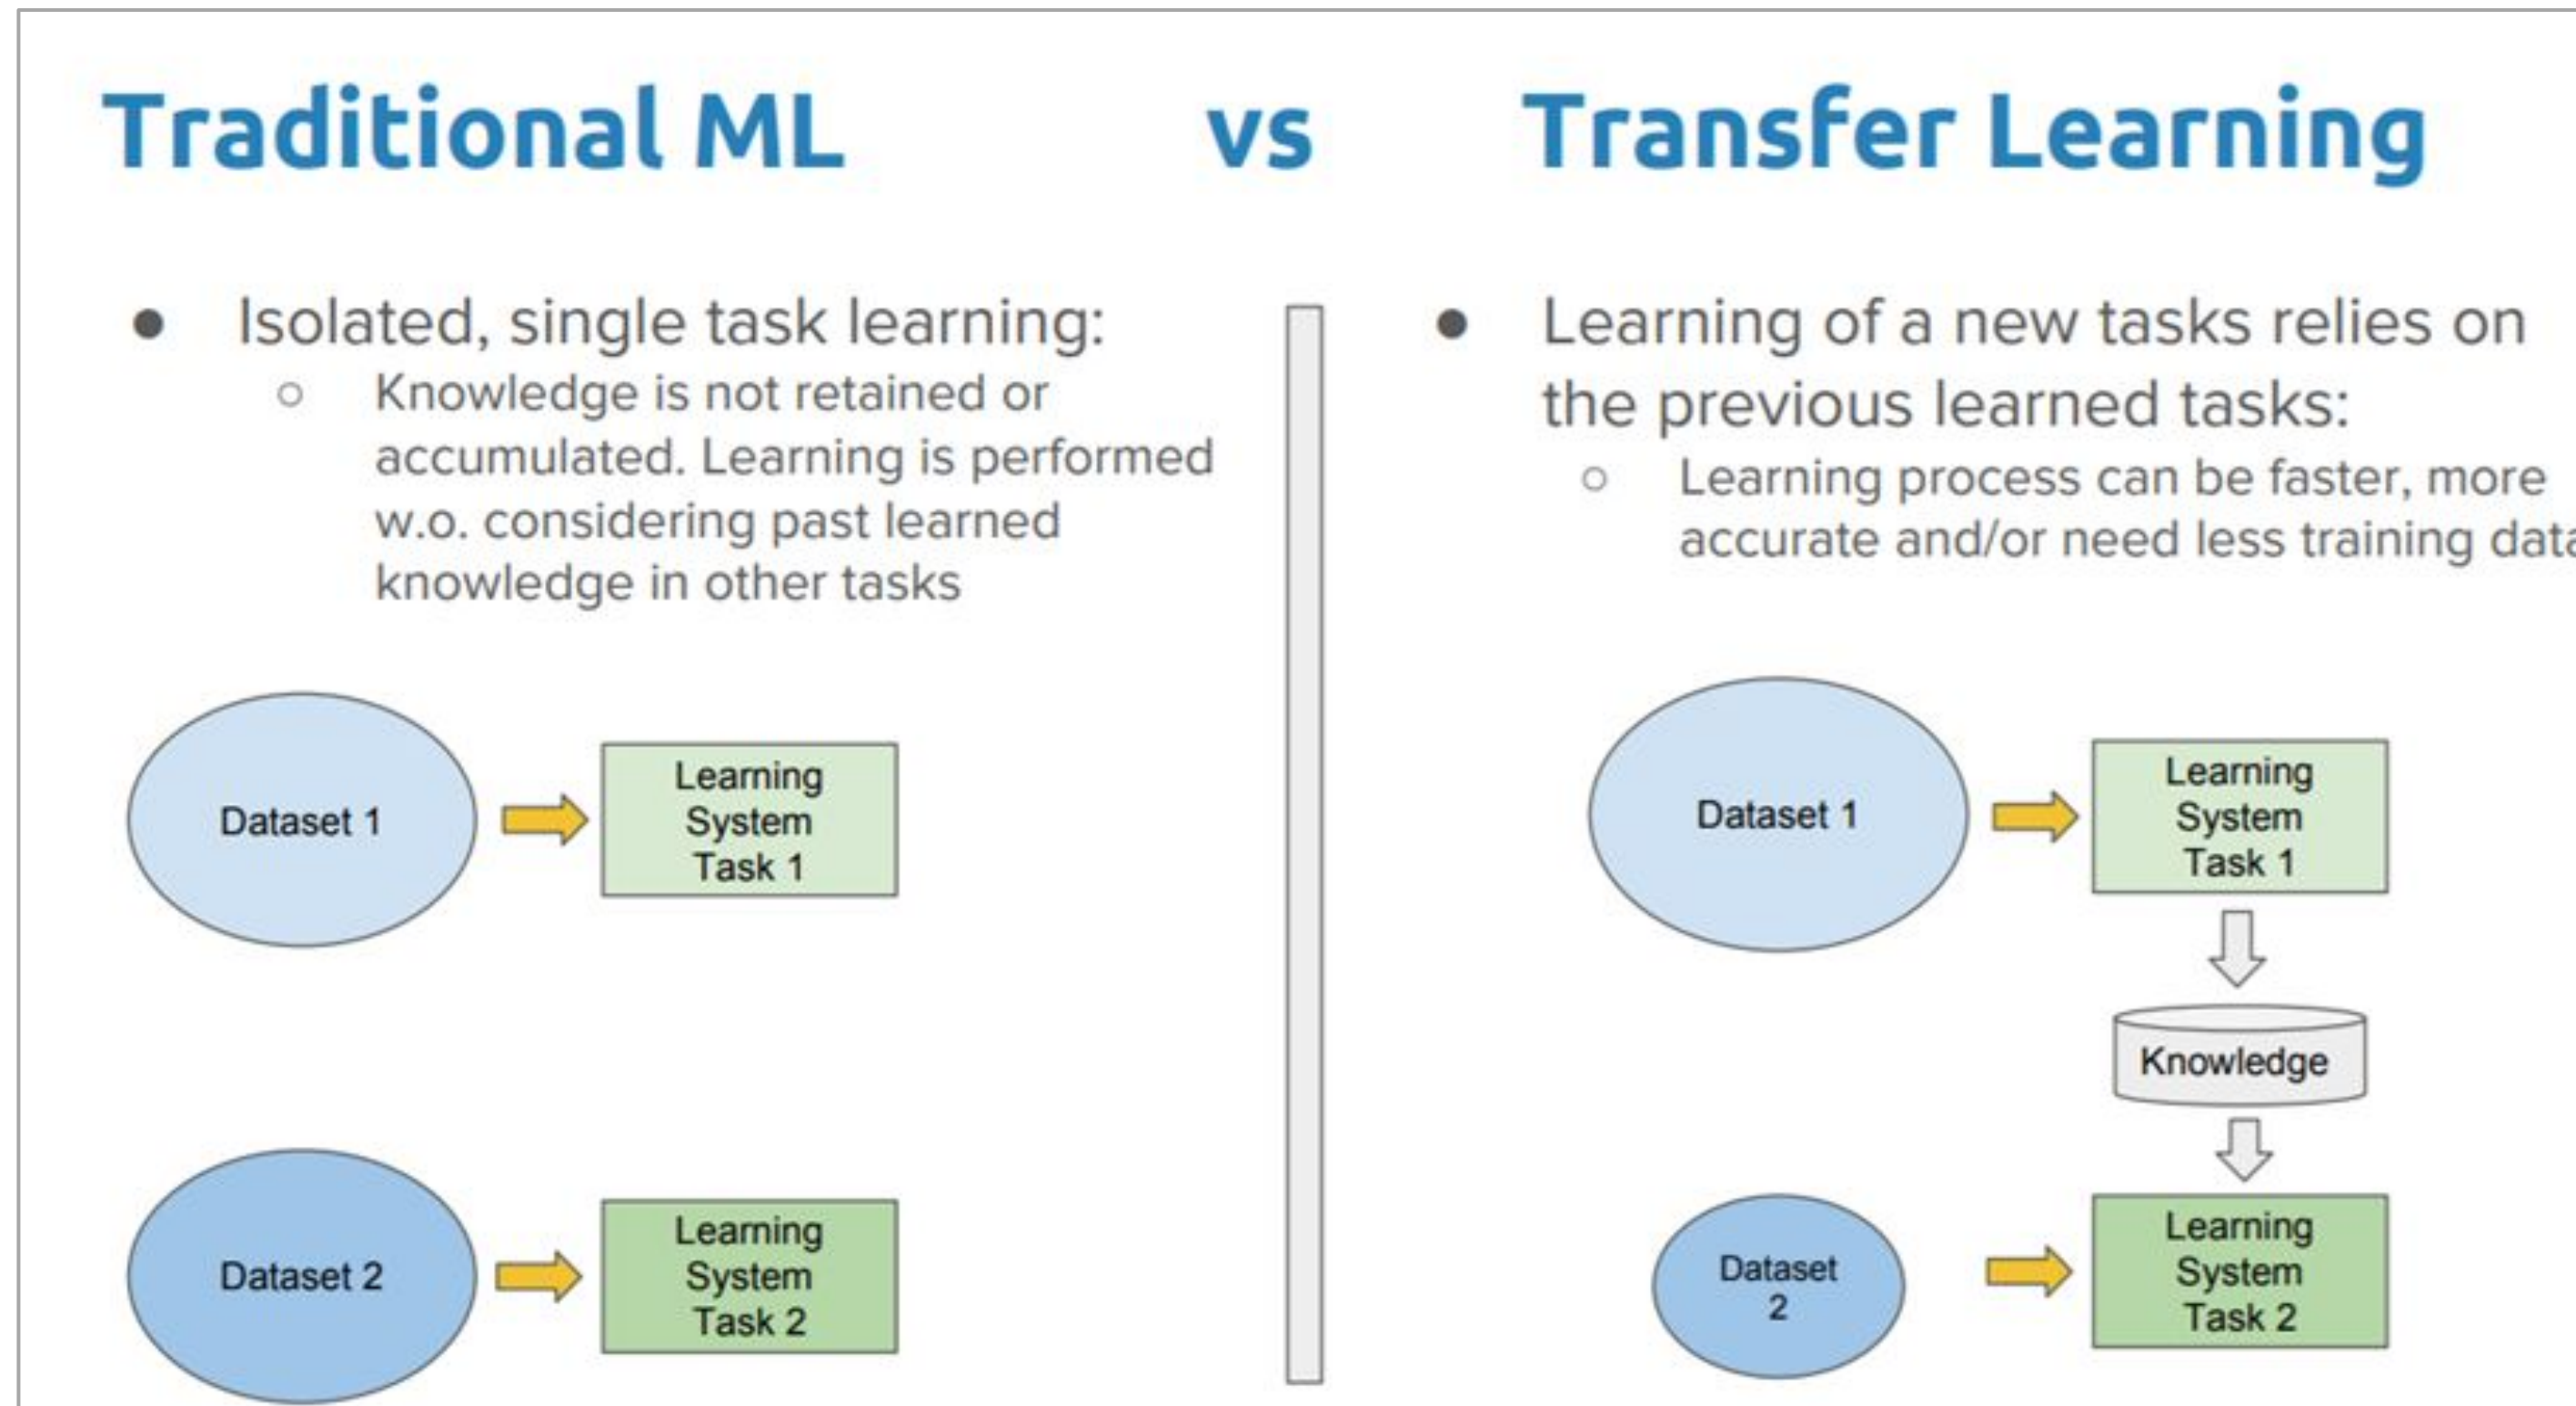
Transfer learning: Intro



Transfer learning: Intro

“After supervised learning — Transfer Learning will be the next driver of ML commercial success.” - Andrew NG ([source](#))

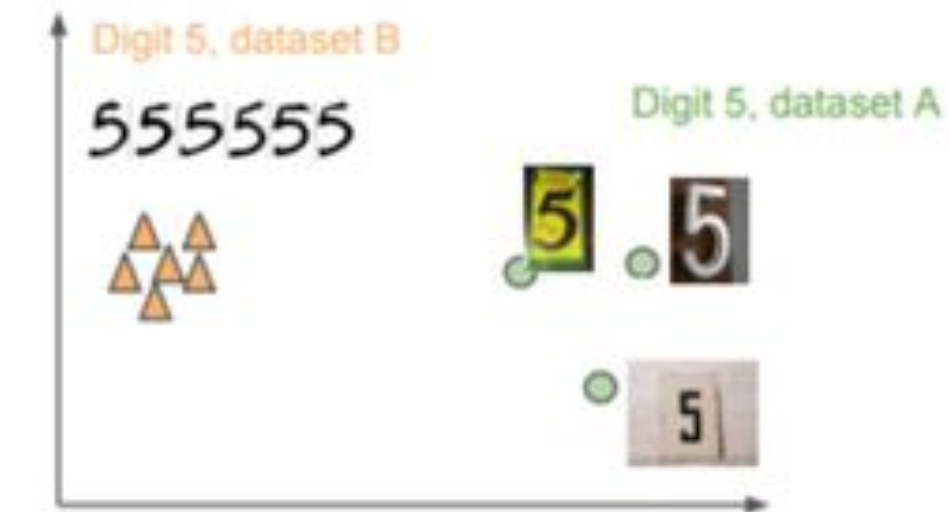
Transfer learning: Intro



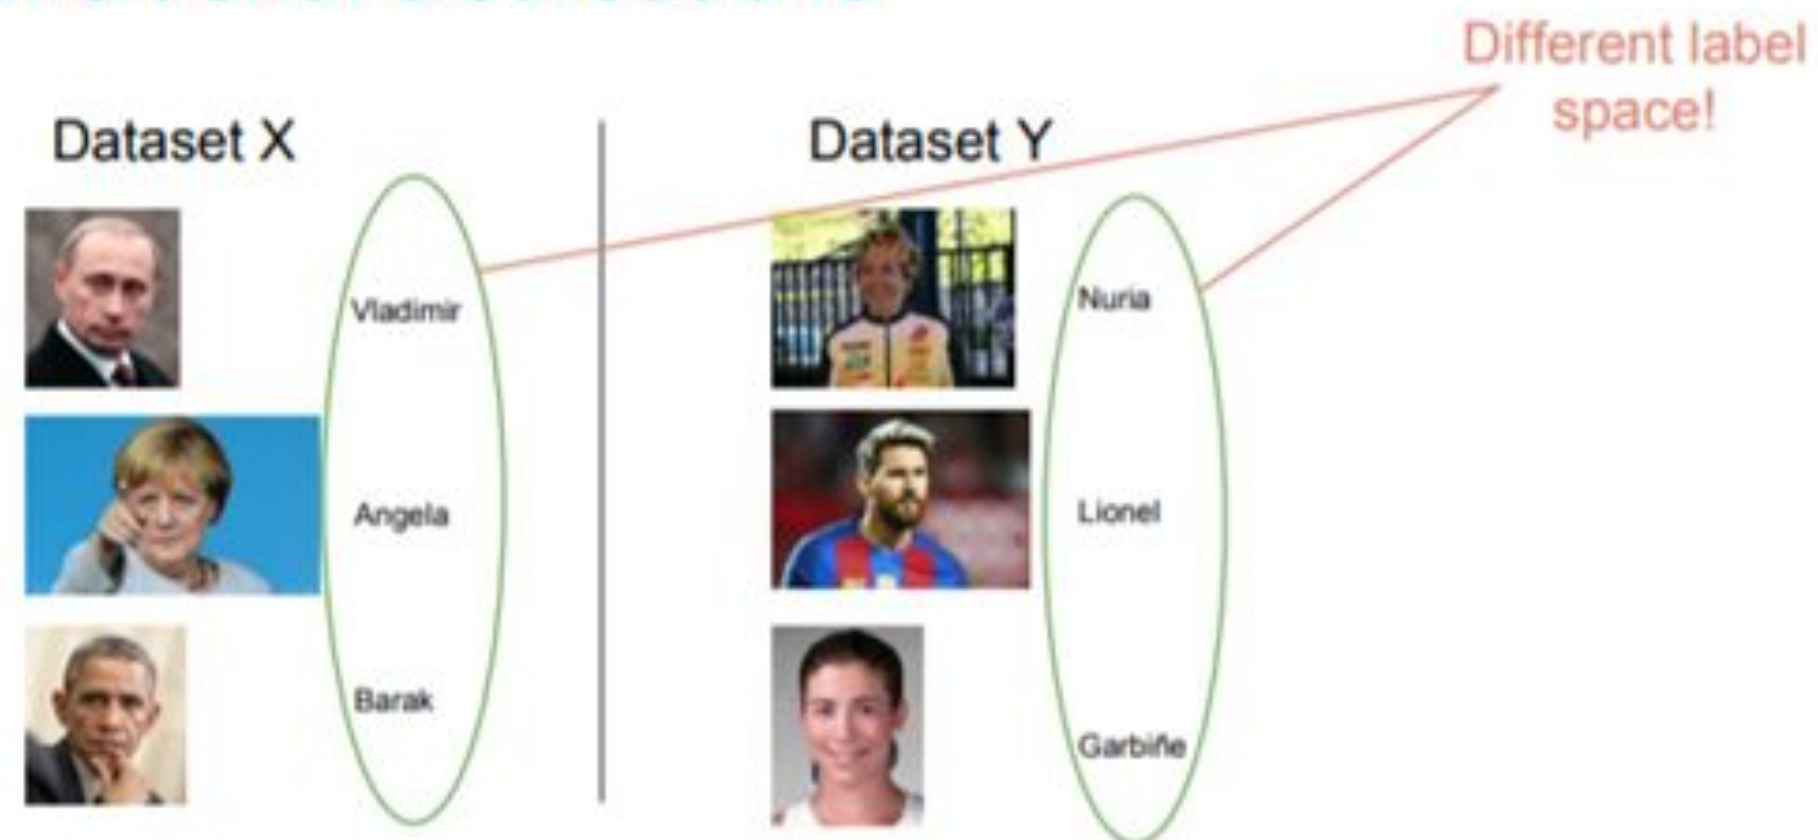
[Source](#)

Transfer learning: Intro

If two domains are different, they may have different **feature spaces** or different **marginal distributions**



If two tasks are different, they may have different **label spaces** or different **conditional distributions**



[Source](#)

Transfer learning: Intro

For a more formal treatment of transfer learning, refer to [Transfer Learning - Machine Learning's Next Frontier](#) by Sebastian Ruder.

The ImageNet moment

ImageNet Challenge

IMAGENET

- 1,000 object classes (categories).
- Images:
 - 1.2 M train
 - 100k test.



[Source](#)

A seminal moment

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

[Source](#)

A seminal moment

- It opened numerous new possibilities which previously seemed impossible.

A seminal moment

- It opened numerous new possibilities which previously seemed impossible.
- Researchers developed models that [achieved superhuman performance](#).

A seminal moment

- It opened numerous new possibilities which previously seemed impossible.
- Researchers developed models that [achieved superhuman performance](#).
- The knowledge learned in those models turned out to be ***gold*** for computer vision tasks.

Representation learning ftw!



[Source](#)

Representation learning ftw!

- Turns out that these low-level features actually constitutes the basic structure of images.

Representation learning ftw!

- Turns out that these low-level features actually constitutes the basic structure of images.
- These low-level features help to develop an understanding of the underlying structures of the images.

Representation learning ftw!

- Turns out that these low-level features actually constitutes the basic structure of images.
- These low-level features help to develop an understanding of the underlying structures of the images.
- These representations are applicable to tasks having visual input modalities -

Representation learning ftw!

- Turns out that these low-level features actually constitutes the basic structure of images.
- These low-level features help to develop an understanding of the underlying structures of the images.
- These representations are applicable to tasks having visual input modalities -
 - Image classification

Representation learning ftw!

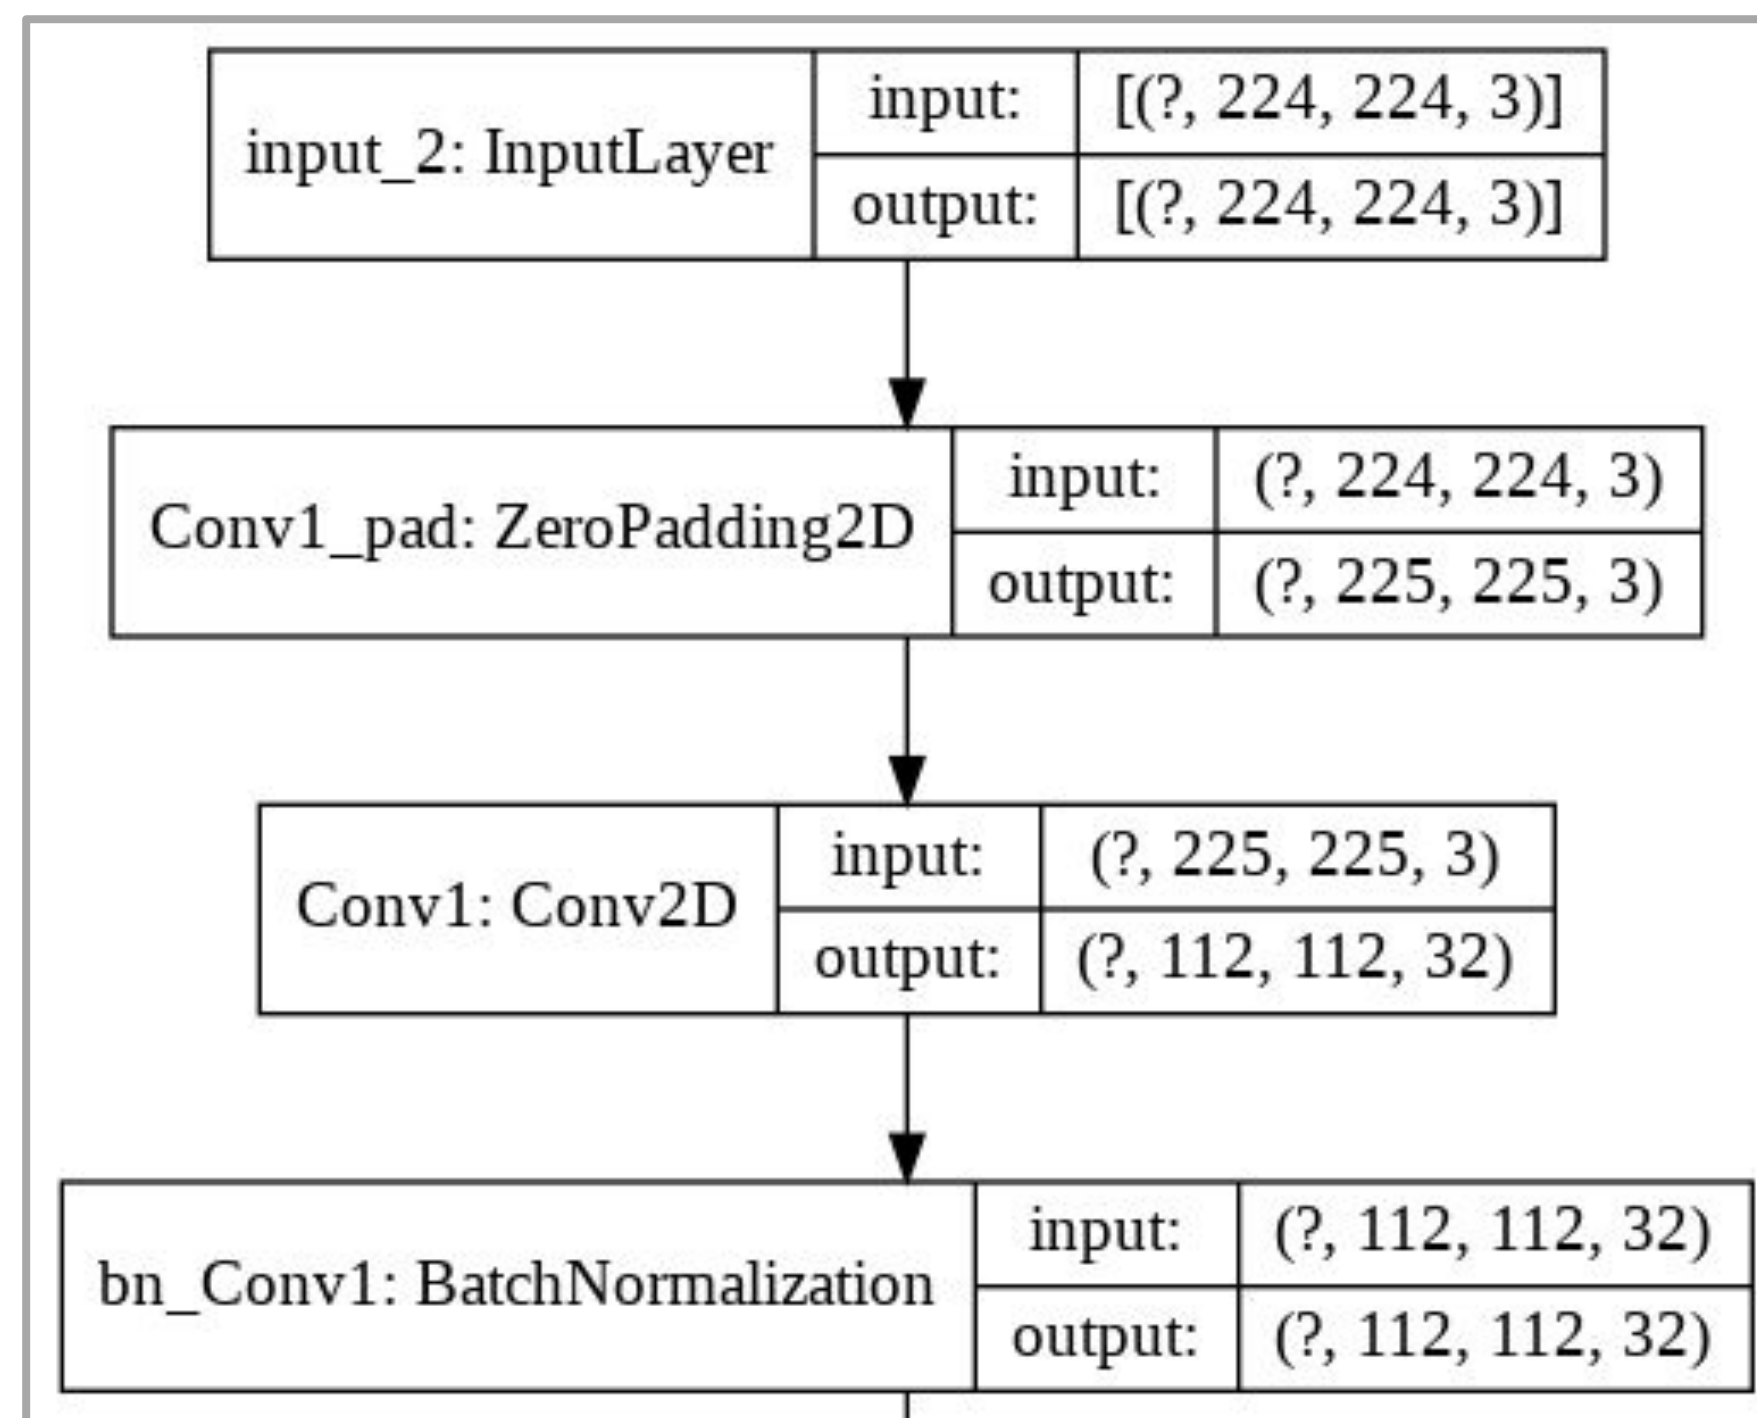
- ...
- These low-level features help to develop an understanding of the underlying structures of the images.
- These representations are applicable to tasks having visual input modalities –
 - Image classification
 - Object detection

Representation learning ftw!

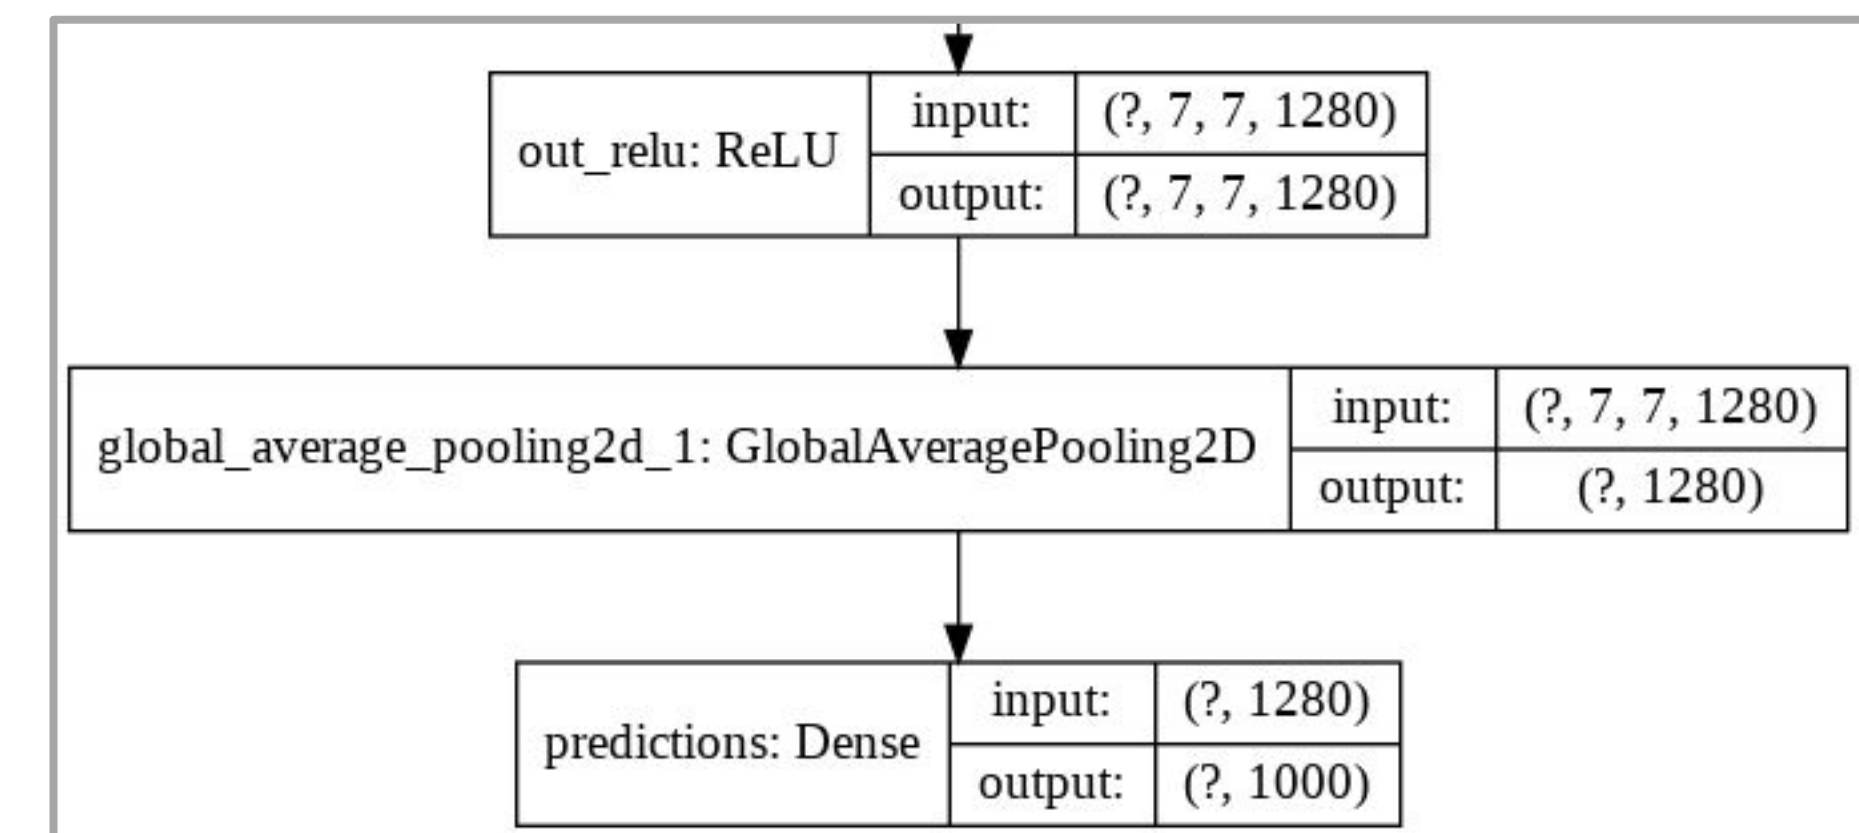
- ...
- These low-level features help to develop an understanding of the underlying structures of the images.
- These representations are applicable to tasks having visual input modalities –
 - Image classification
 - Object detection
 - Image captioning
 - and [many more!](#)

Transfer learning for image classification

A sample architecture (MobileNetV2) -



First half of the network




Classification top of the network

Transfer learning for image classification

Three recipes -

- Off-the-shelf inference.
- Precomputing bottlenecks and replacing the classification top.
- Fine-tuning.

Off-the-shelf inference



```
# Load off the base model
```


```
mobilenet = MobileNetV2(weights="imagenet")
```

```
# Run inference and parse the predictions
```

```
result = mobilenet.predict(image)
```

```
print(decode_predictions(result, top=3)[0])
```


Precomputing *bottlenecks* and replacing the classification top

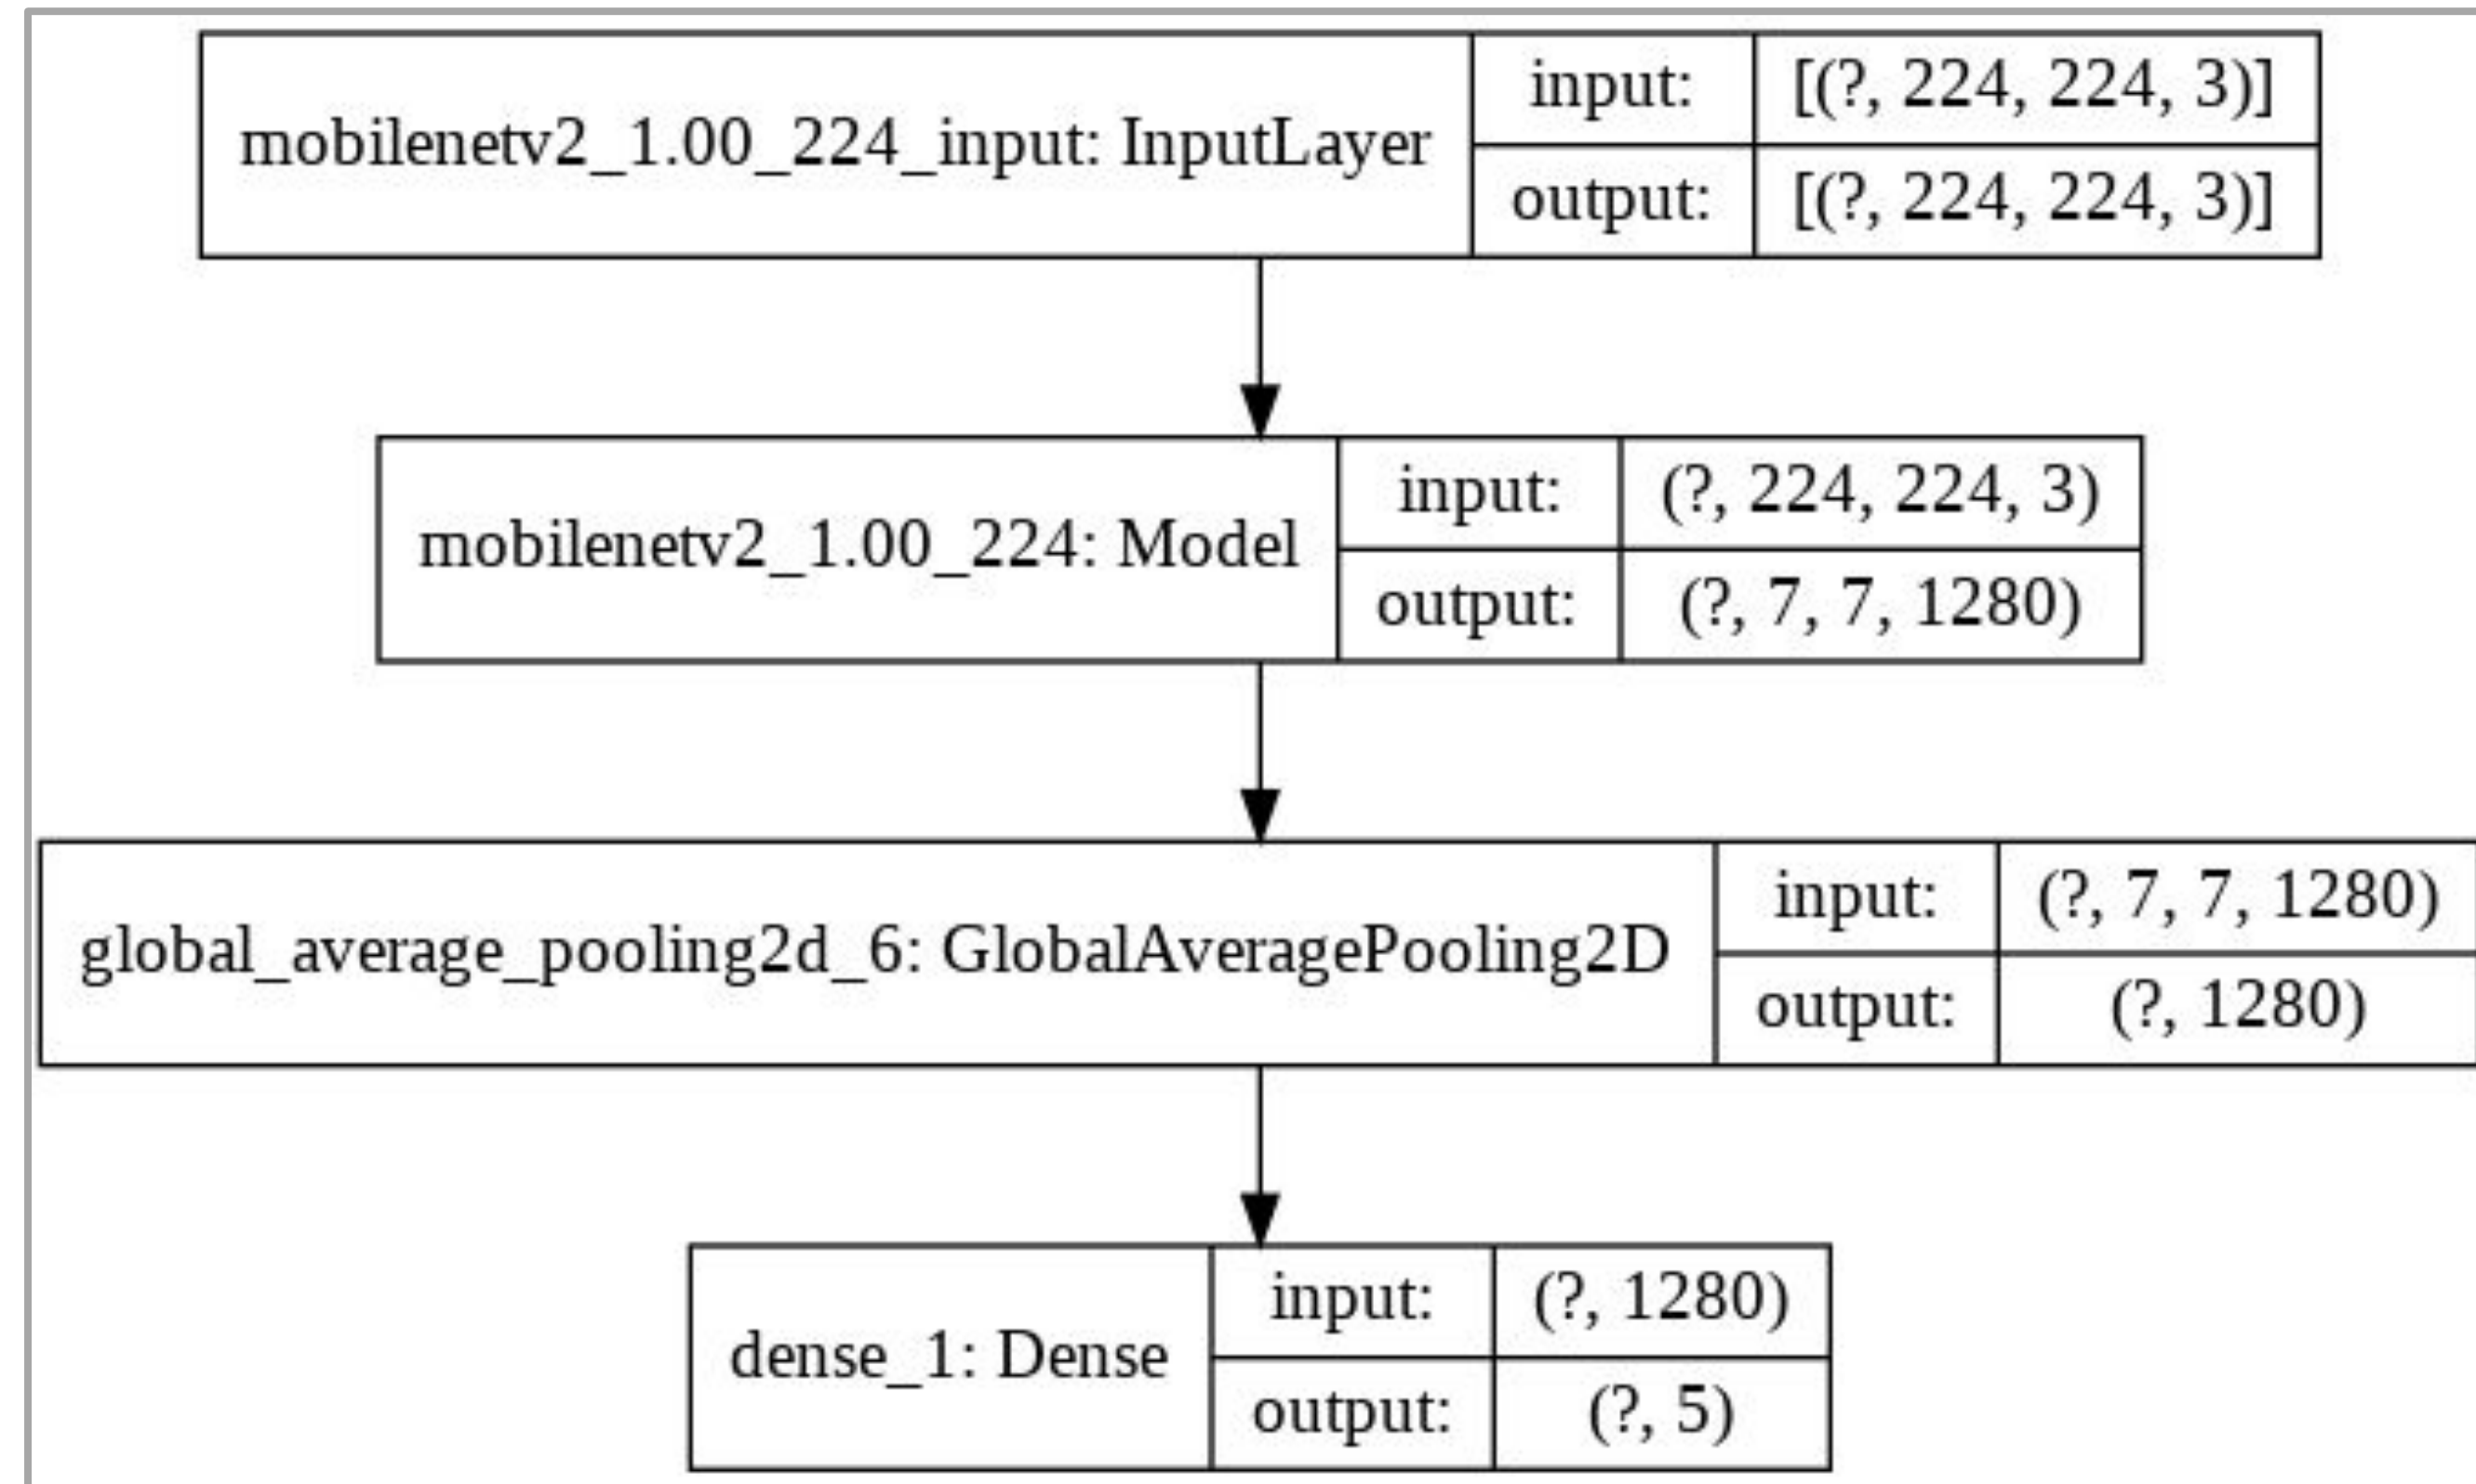


```
# Load the pre-trained model without classification top
mobilenet = MobileNetV2(weights="imagenet", include_top=False)

# Set the base model to non-trainable
mobilenet.trainable = False

# Build the new model
new_model = Sequential([
    mobilenet,
    GlobalAveragePooling2D(),
    Dense(len(CLASSES), activation="softmax")
])
```

Resultant model



Resultant model

```
1 new_model.summary()
```


Model: "sequential_1"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Model)	(None, 7, 7, 1280)	2257984
global_average_pooling2d_6 ((None, 1280)	0
dense_1 (Dense)	(None, 5)	6405

Total params: 2,264,389
Trainable params: 6,405
Non-trainable params: 2,257,984

Number of trainable parameters very low as expected

Fine-tuning




```
# Load the pre-trained model without classification top
mobilenet = MobileNetV2(weights="imagenet", include_top=False)

# Set the base model to *trainable*
mobilenet.trainable = True

# Build the new model
new_model_trainable = Sequential([
    mobilenet,
    GlobalAveragePooling2D(),
    Dense(len(CLASSES), activation="softmax")
])
```

Resultant model

```
1 new_model_trainable.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Model)	(None, 7, 7, 1280)	2257984
global_average_pooling2d_7 ((None, 1280)	0
dense_2 (Dense)	(None, 5)	6405

Total params: 2,264,389
Trainable params: 2,230,277
Non-trainable params: 34,112

Number of trainable parameters
changes as expected

Fine-tuning

Can be done in two variants:

- After doing a round plain transfer learning.
- Training the pre-trained network from the beginning with **trainable=True**.

Things to keep in mind

- When pre-computing the bottlenecks/embeddings:
 - Be careful about allocating resources.

Things to keep in mind

- When pre-computing the bottlenecks/embeddings:
 - Be careful about allocating resources.

```
1 new_model.summary()

Model: "sequential_1"

Layer (type)                 Output Shape              Param #
=====
mobilenetv2_1.00_224 (Model) (None, 7, 7, 1280)        2257984
global_average_pooling2d_6 ( (None, 1280)          0
dense_1 (Dense)               (None, 5)                 6405
=====
Total params: 2,264,389
Trainable params: 6,405
Non-trainable params: 2,257,984
```

You are only training a ***shallow fully-connected*** network.

Things to keep in mind

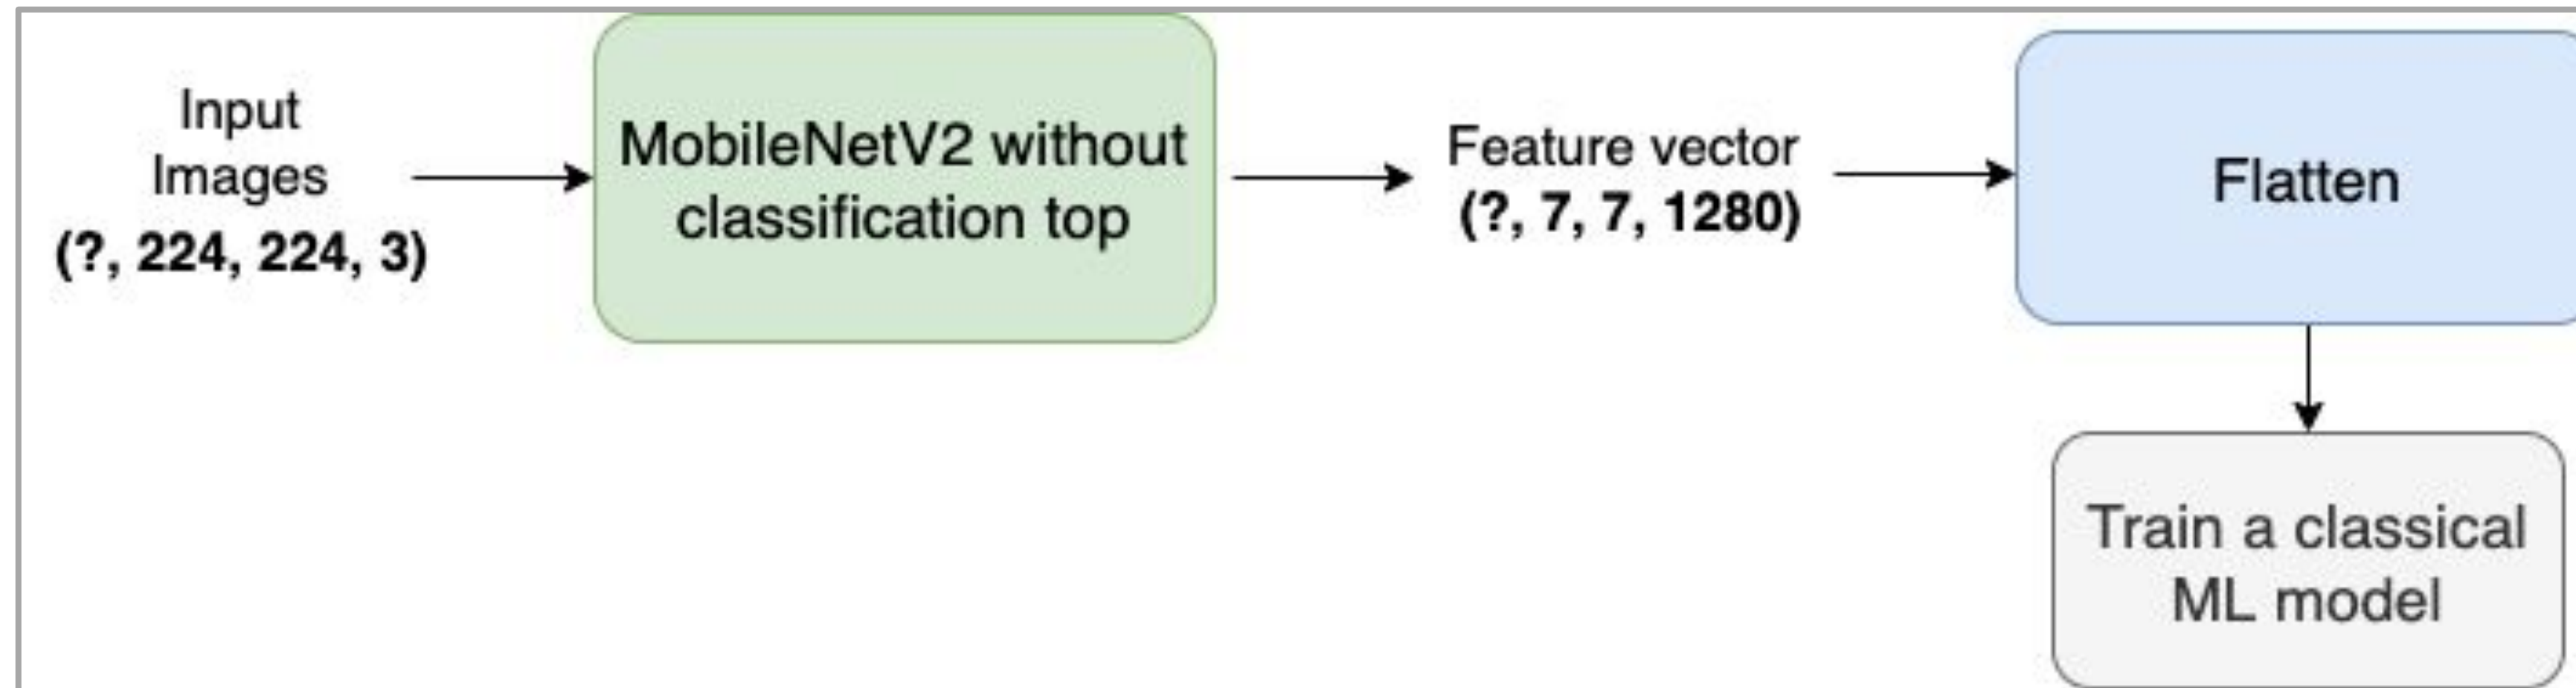
- When pre-computing the bottlenecks/embeddings:
 - Be careful about allocating resources.
 - Try with a ***relatively higher learning rate*** in this case.

Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.

Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.



An example [here](#)

Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.
- When doing fine-tuning:
 - Start with a ***lower learning rate*** -

Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.
- When doing fine-tuning:
 - Start with a ***lower learning rate*** -
 - Discriminative learning rates
 - Learning rate schedules

Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.
- When doing fine-tuning?
- Mean subtraction with the source dataset mean stats.

```
# Initialize your data generator
```

```
train_aug = ImageDataGenerator()
```

```
# ImageNet (source dataset) mean stats and assign
```

```
mean = np.array([123.68, 116.779, 103.939], dtype="float32")
```

```
train_aug.mean = mean
```

Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.
- When doing fine-tuning?
- Mean subtraction with the source dataset mean stats.
- Set the batchnorm layers to be ***non-trainable***.

```
# Freeze the base_model
```

```
base_model.trainable = False
```

```
# Create new model on top
```

```
inputs = keras.Input(shape=(150, 150, 3))
```

```
# Make sure the base model runs in *inference mode*
```

```
x = base_model(x, training=False)
```

```
# Classification top
```

```
x = keras.layers.GlobalAveragePooling2D()(x)
```

```
x = keras.layers.Dropout(0.2)(x)
```

```
outputs = keras.layers.Dense(1)(x)
```

```
model = keras.Model(inputs, outputs)
```


Things to keep in mind

- Pre-computing the bottlenecks/embeddings.
- Pre-trained networks + Classical ML.
- When doing fine-tuning?
- Mean subtraction with the source dataset mean stats.
- Set the batchnorm layers to be ***non-trainable***.
- Gradual unfreezing of layers while doing fine-tuning.
(Should be performed after a round of transfer learning)

```
# Unfreeze some of the upper layers of base model
```

```
for layer in base_model.layers[15:]:
```

```
    layer.trainable = True
```

```
# Compile again and train
```

```
model.compile(...)
```

```
model.fit(...)
```

← **Very important!**

New frontier: Self-supervised learning

- ImageNet is a labeled dataset.

New frontier: Self-supervised learning

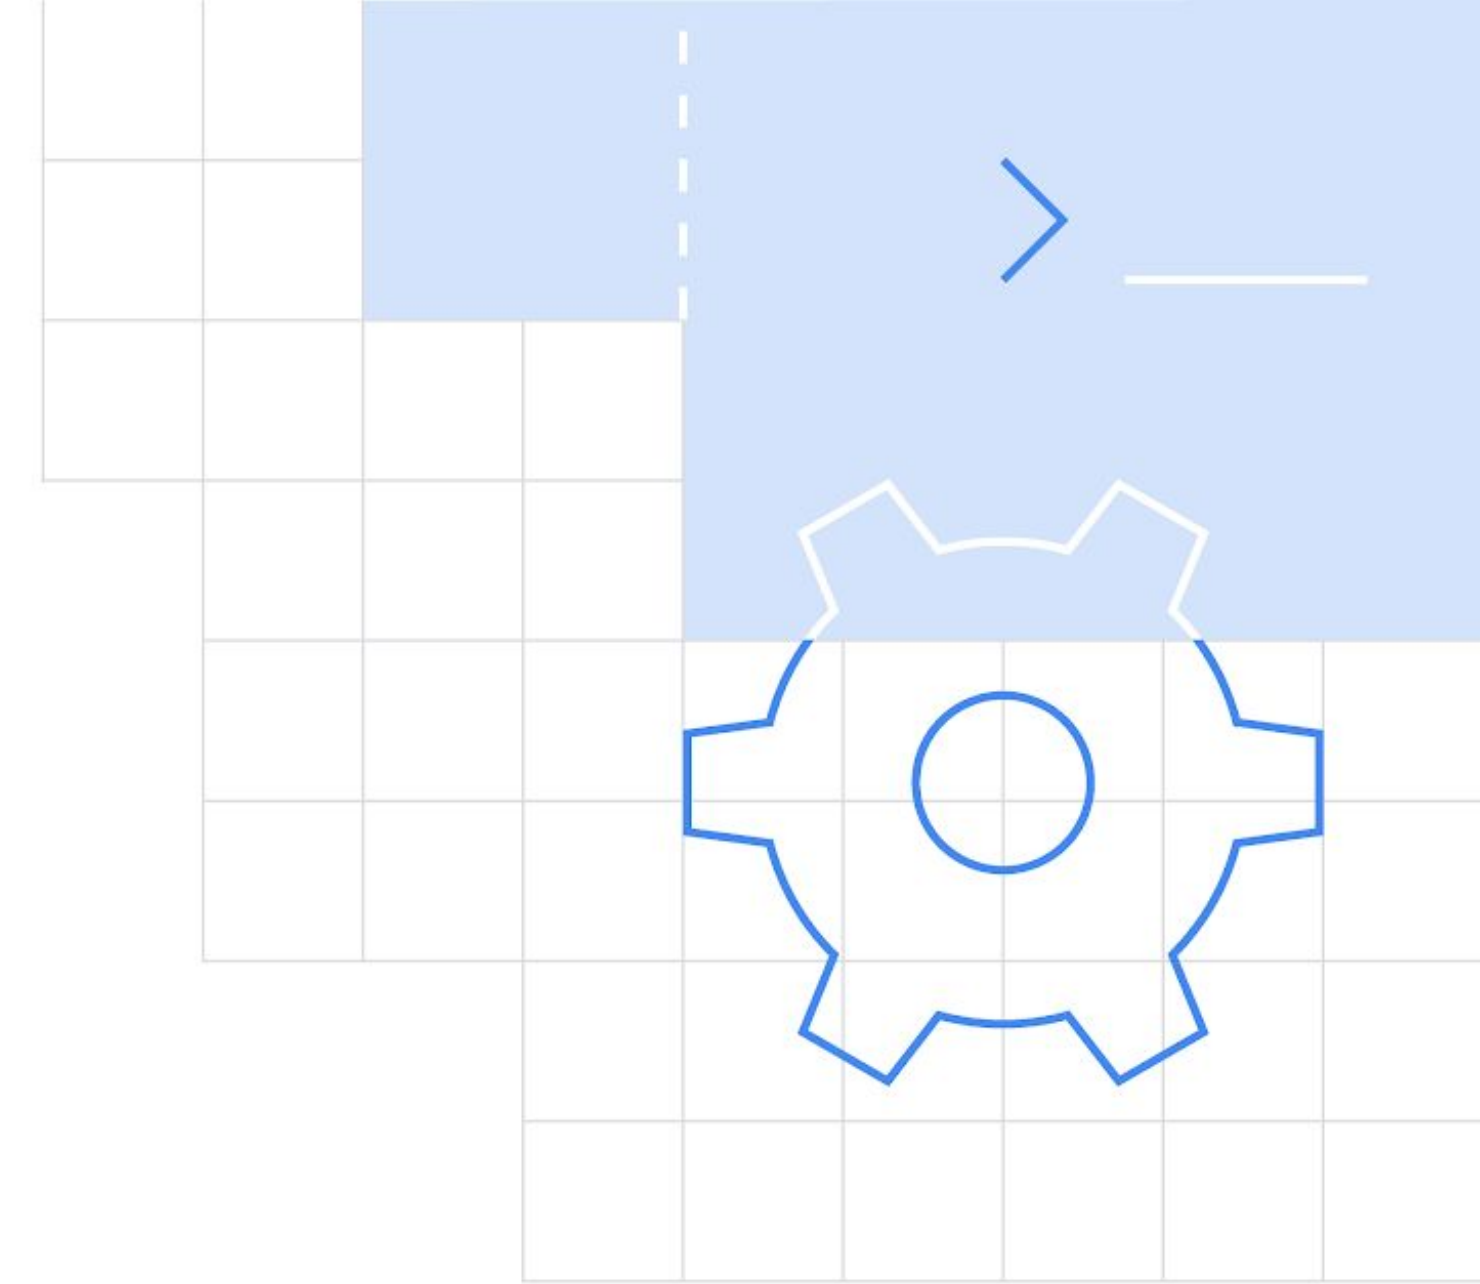
- ImageNet is a labeled dataset.
- Affording such a huge labeled dataset is often not a feasible option for many industries (for example - medical).

New frontier: Self-supervised learning

- ImageNet is a labeled dataset.
- Affording such huge amount of labeled dataset is often not a feasible option for many industries (for example - medical).
- Unlabeled data, whereas, is way cheaper.

New frontier: Self-supervised learning

Check out [**SimCLR**](#) that presents a simple yet effective framework to utilize unlabeled data to learn useful representations from visual data.



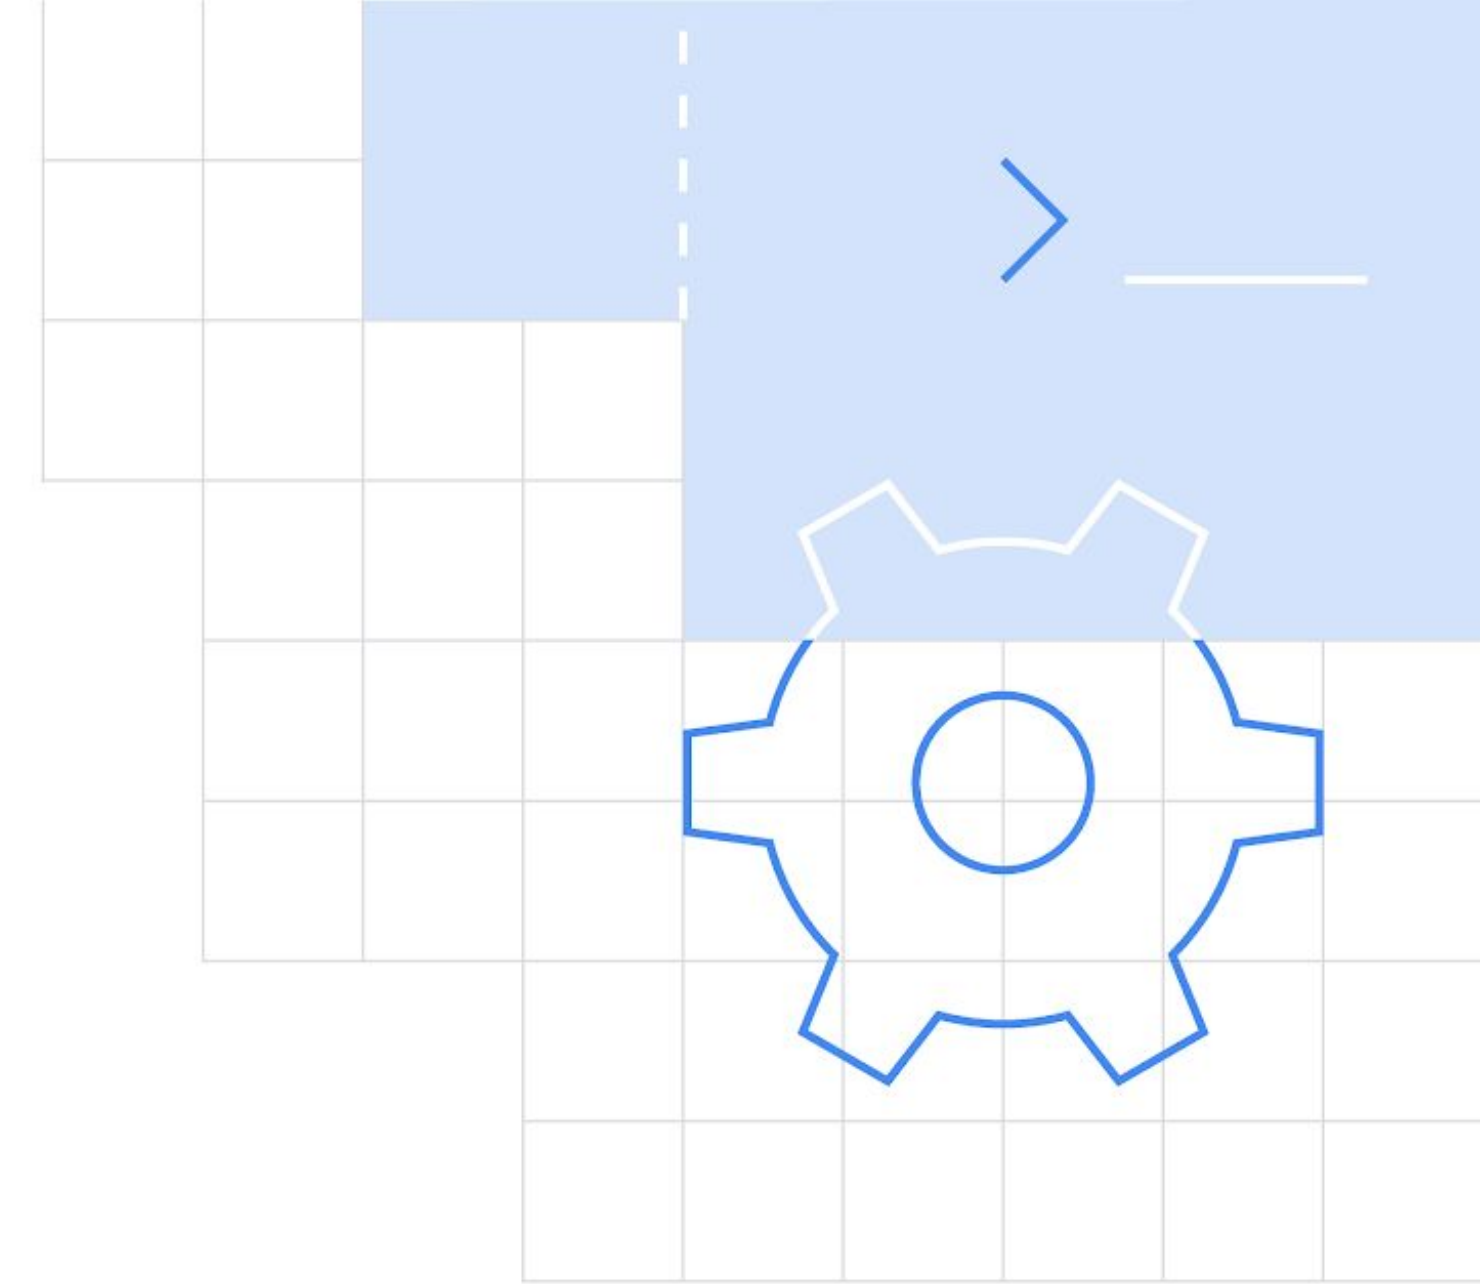
Questions?

References

- [A Comprehensive Hands-on Guide to Transfer Learning with Real-World Applications in Deep Learning](#) by Dipanjan Sarkar
- [Transfer Learning - Machine Learning's Next Frontier](#) by Sebastian Ruder
- [Fine-tuning with Keras and Deep Learning](#) by PyImageSearch
- [Transfer learning notebook](#) by François Chollet

Slides available here -

<https://bit.ly/tl-sayak>



Thank You!



Sayak Paul
PyImageSearch
[@RisingSayak](#)

